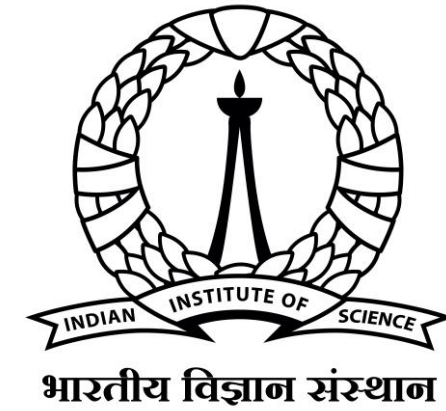




Department of Computational and Data Sciences



Module 4: Generative AI

Deepak Subramani

Assistant Professor

Dept. of Computational and Data Science

Indian Institute of Science Bengaluru

Outline for Jan 25

- Part 1: Decoder only GPT Model
 - What are GPT-class Generative Large Language Models
 - Data preparation for GPT model training
 - GPT finetuning (Assignment)
- Part 2: LLMs and Interacting with them
 - Commercial and open source LLMs
 - What are the main issues in LLMs to be aware of?
 - Taxonomy of interaction with LLMs
 - Prompting Strategies – ZSL, FSL, CoT, ReACT, DSP
 - Parameter Efficient Fine Tuning (LoRA, QLoRA)



Department of Computational and Data Sciences

Outline for Feb 01



- Part 1:
 - Instruction Tuning
- Part 2:
 - Orchestration
 - Retrieval Augmented Generation
- Part 3:
 - LLM Guardrails
 - LLM Agents



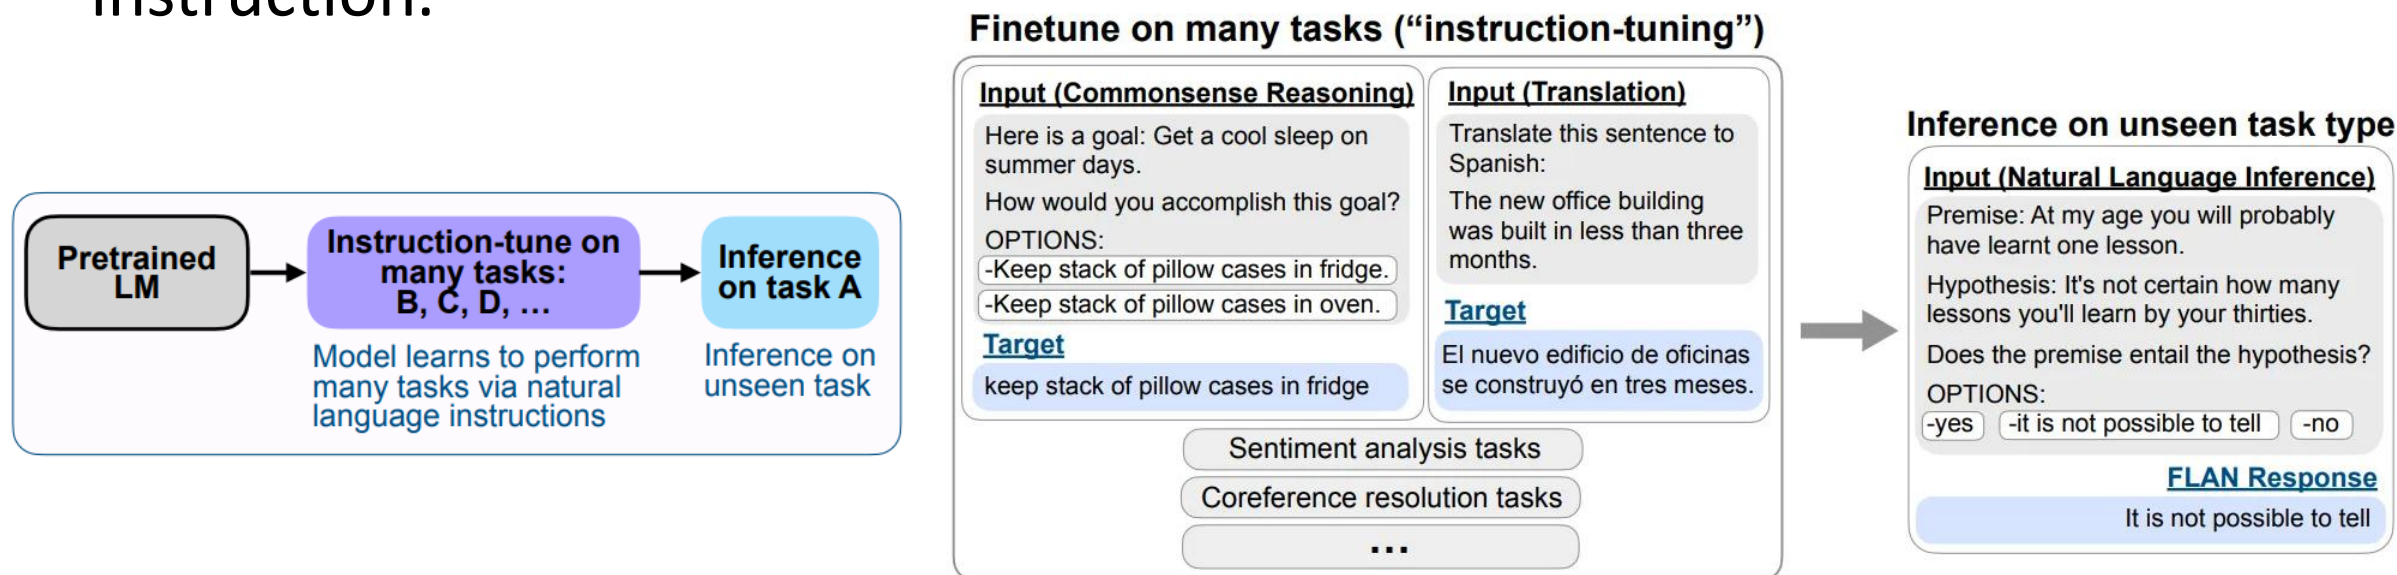
List of Prompting Strategies

1. Chain of Thought Prompting (CoT)
2. Tree of Thought Prompting
3. Prompts with Instructions
4. Prompt Chaining
5. Directed Stimulus Prompting
6. ReAct



Prompts with Instructions

- Unlike open-ended prompts, prompts with instructions guide the model to perform specific tasks or follow certain rules within its response.
- **Structured Responses:** Instruction-based prompts lead to more structured and targeted outputs, adhering to the format or content specified in the instruction.






Instruct Fine Tuning via RLHF


Step 1

Collect demonstration data, and train a supervised policy.


A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.




This data is used to fine-tune GPT-3 with supervised learning.




Step 2

Collect comparison data, and train a reward model.


A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.



This data is used to train our reward model.




Step 3

Optimize a policy against the reward model using reinforcement learning.

A new prompt is sampled from the dataset.



The policy generates an output.




Once upon a time...

The reward model calculates a reward for the output.



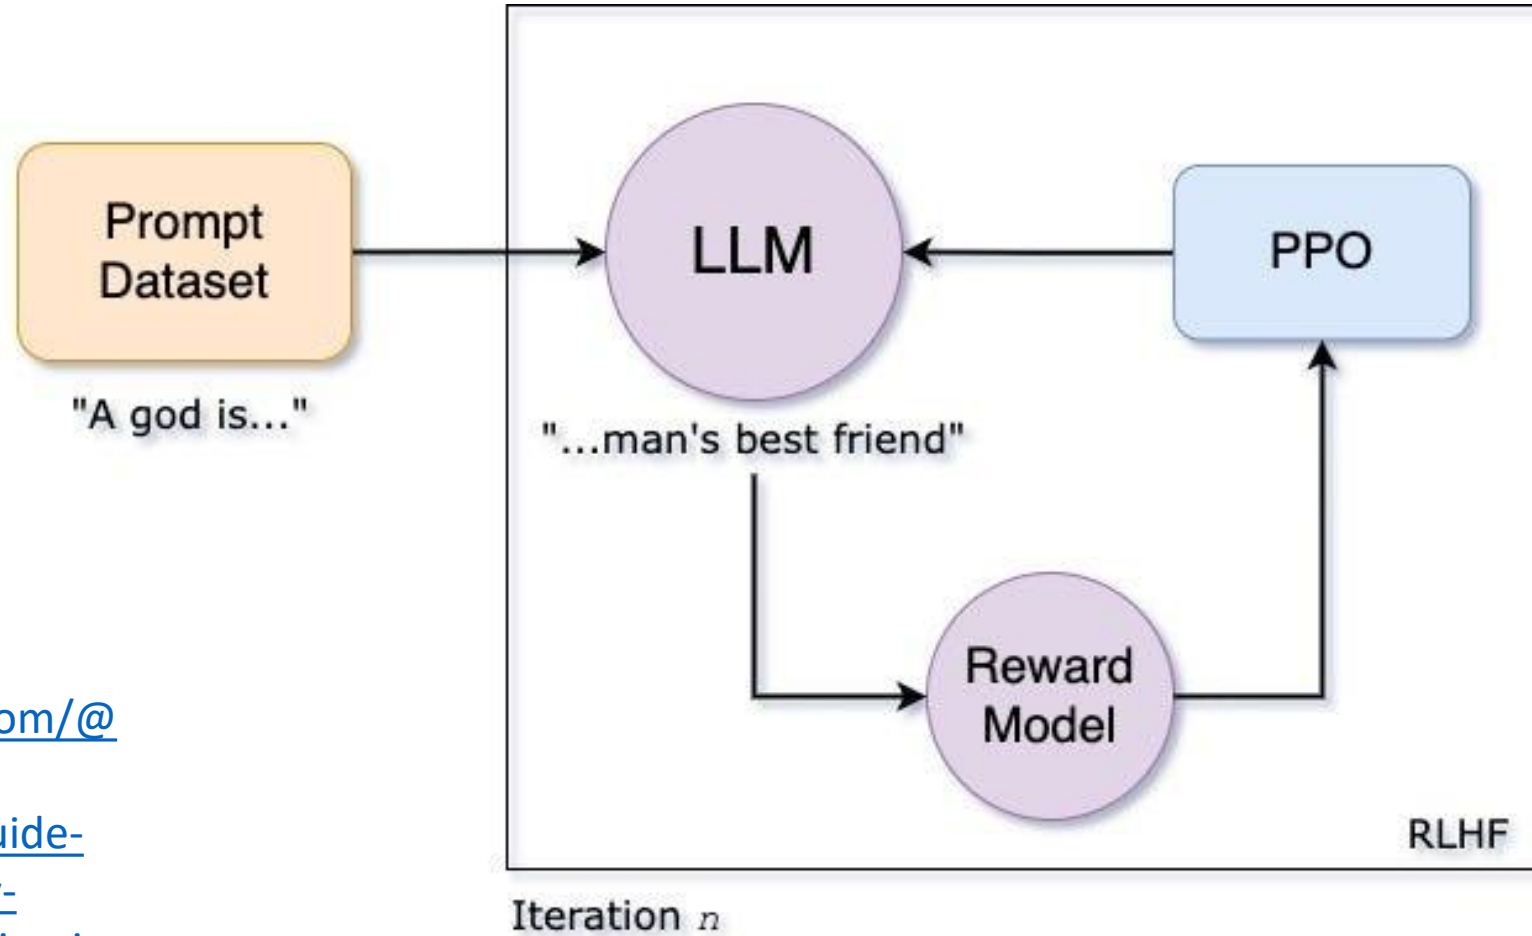
The reward is used to update the policy using PPO.



r_k



RLHF via Proximal Policy Optimization (PPO)

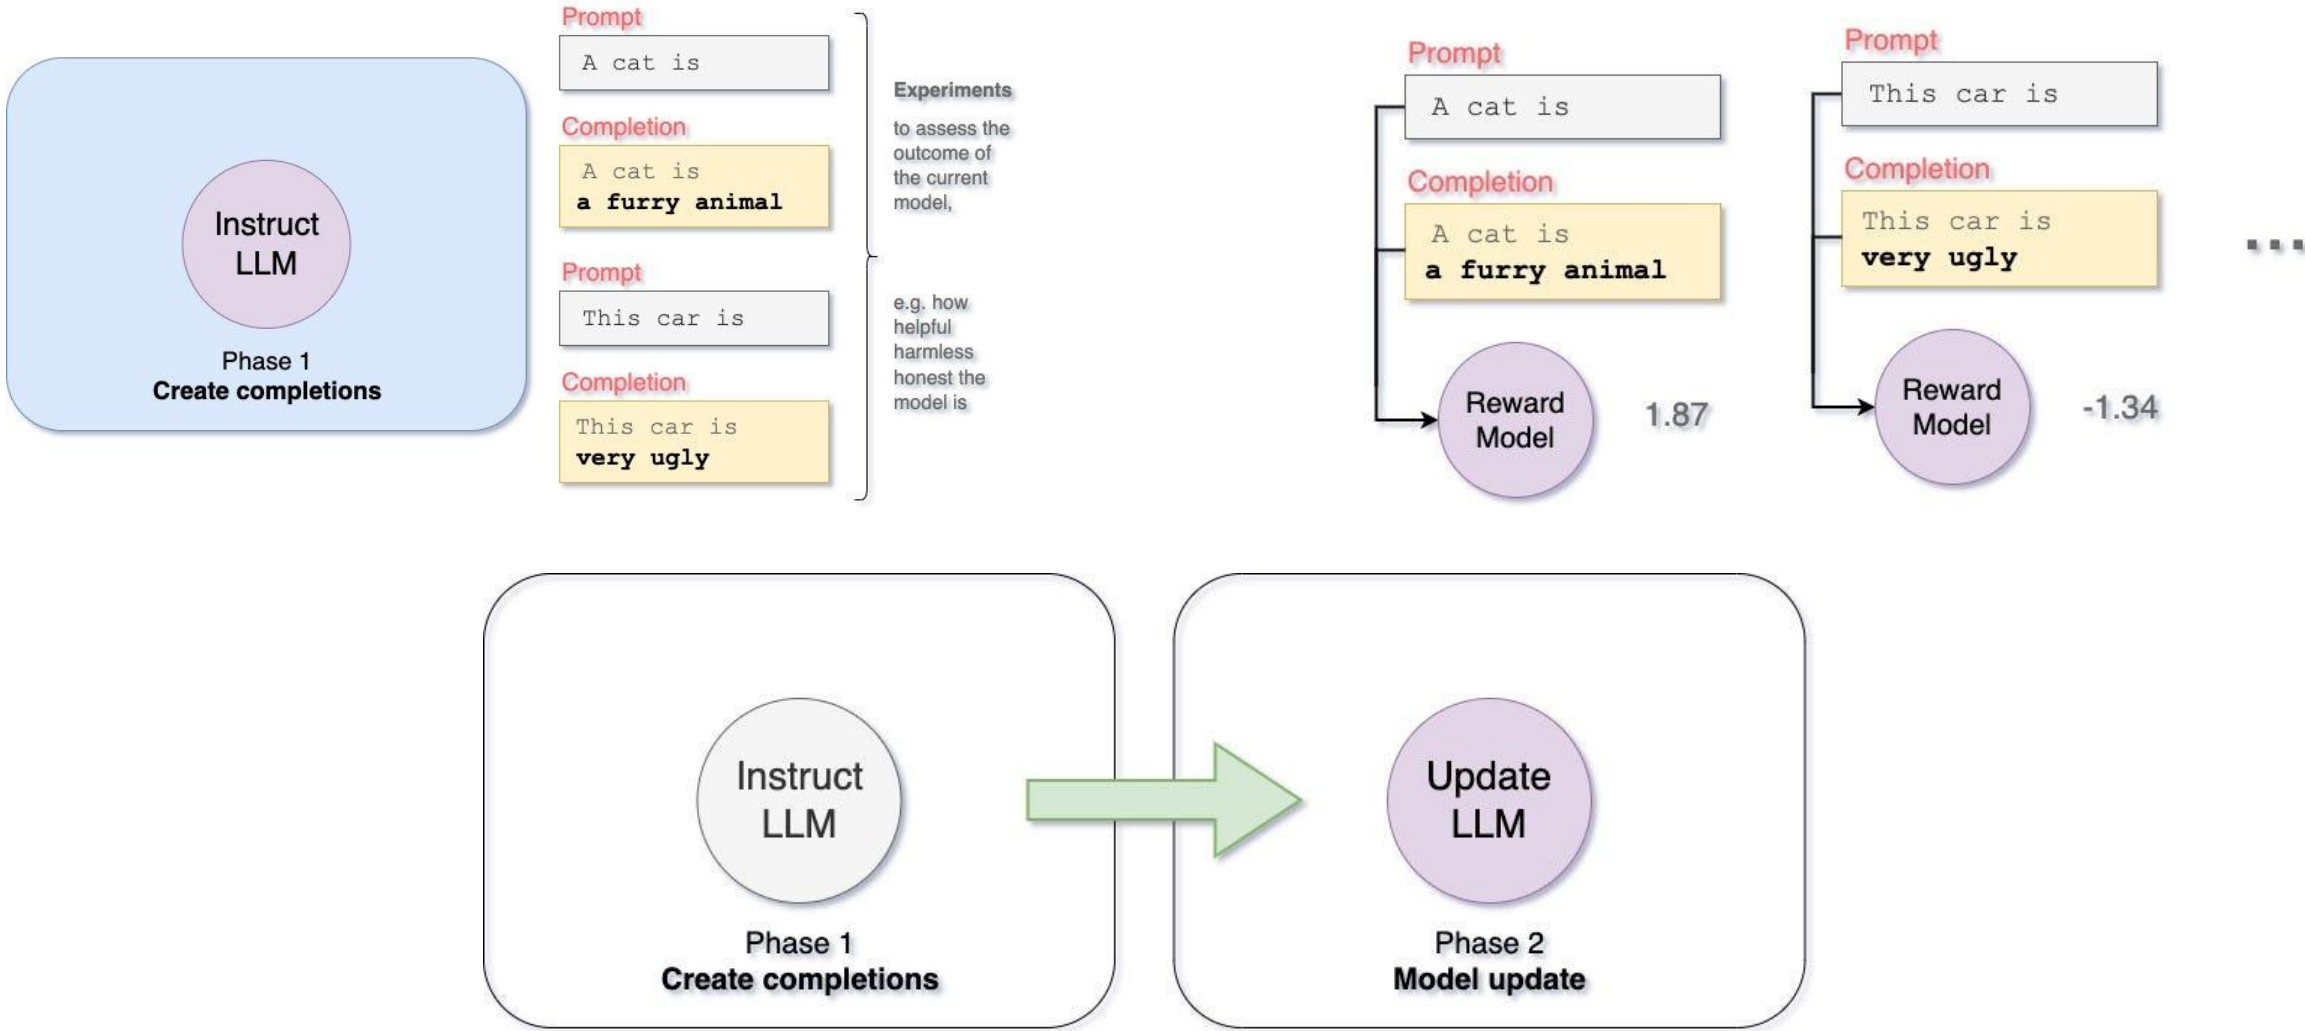


Source:

<https://medium.com/@oleglatypov/a-comprehensive-guide-to-proximal-policy-optimization-ppo-in-ai-82edab5db200>



PPO Steps





The RL Setting

- 5 Key Concepts

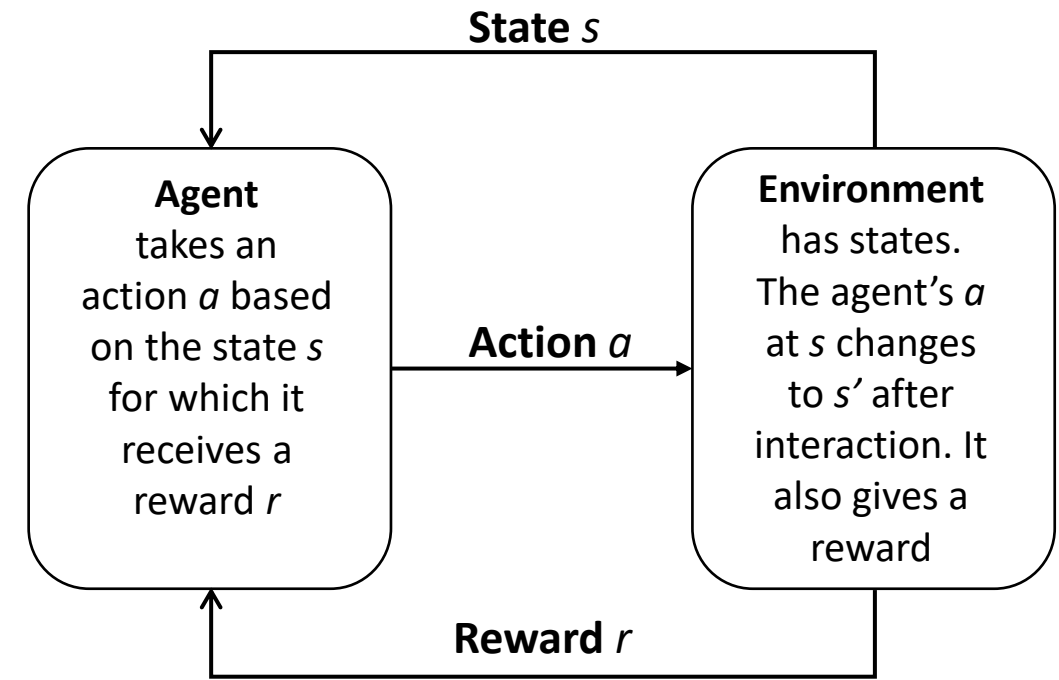

Agent

Environment

State (Observation)

Action

Reward





Department of Computational and Data Sciences

Outline for Feb 01



- Part 1:
 - Instruction Tuning
- Part 2:
 - Orchestration
 - Retrieval Augmented Generation
- Part 3:
 - LLM Guardrails
 - LLM Agents



LangChain

Data connection

Document
loaders

Document
transformers

Embedding
models

Vector
stores

Prompts



LangChain.dart

Retrievers

Language
models

Agents

Output
parsers

Chains

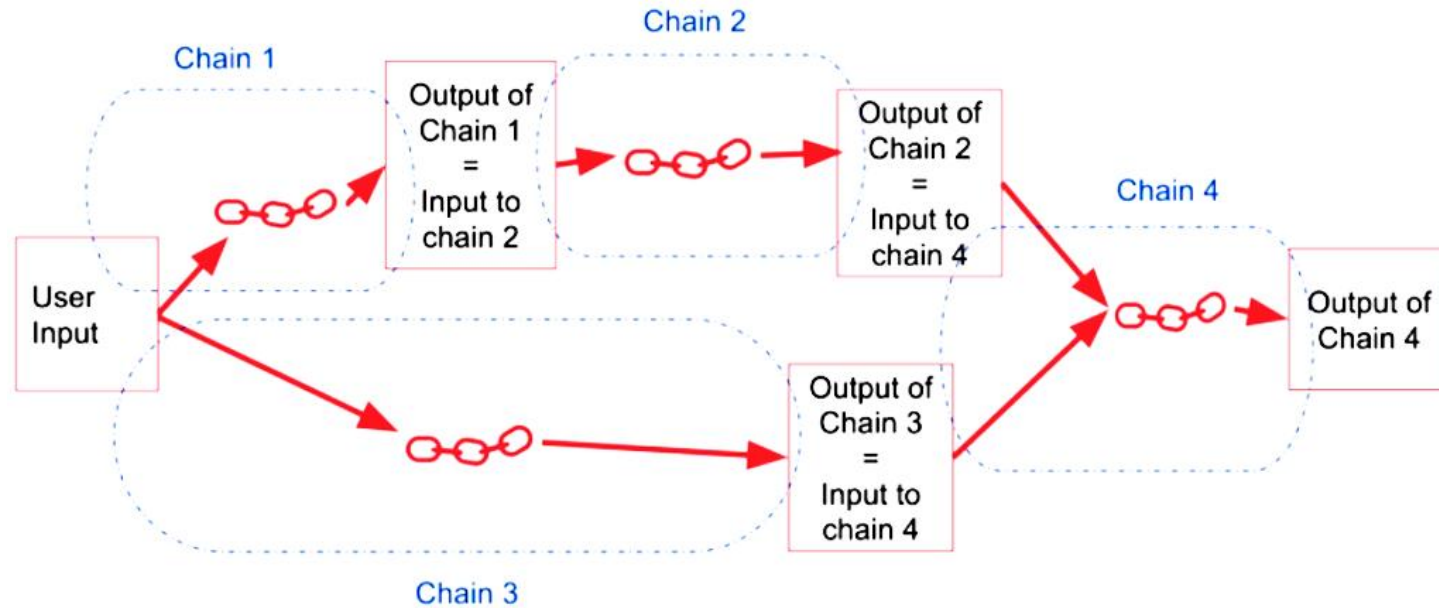
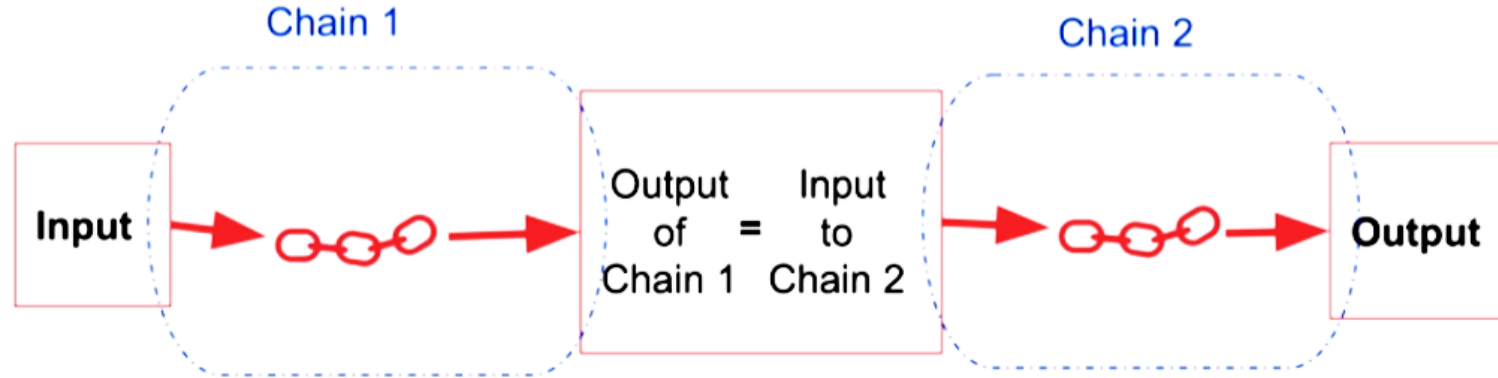
Memory

Tools

Model I/O

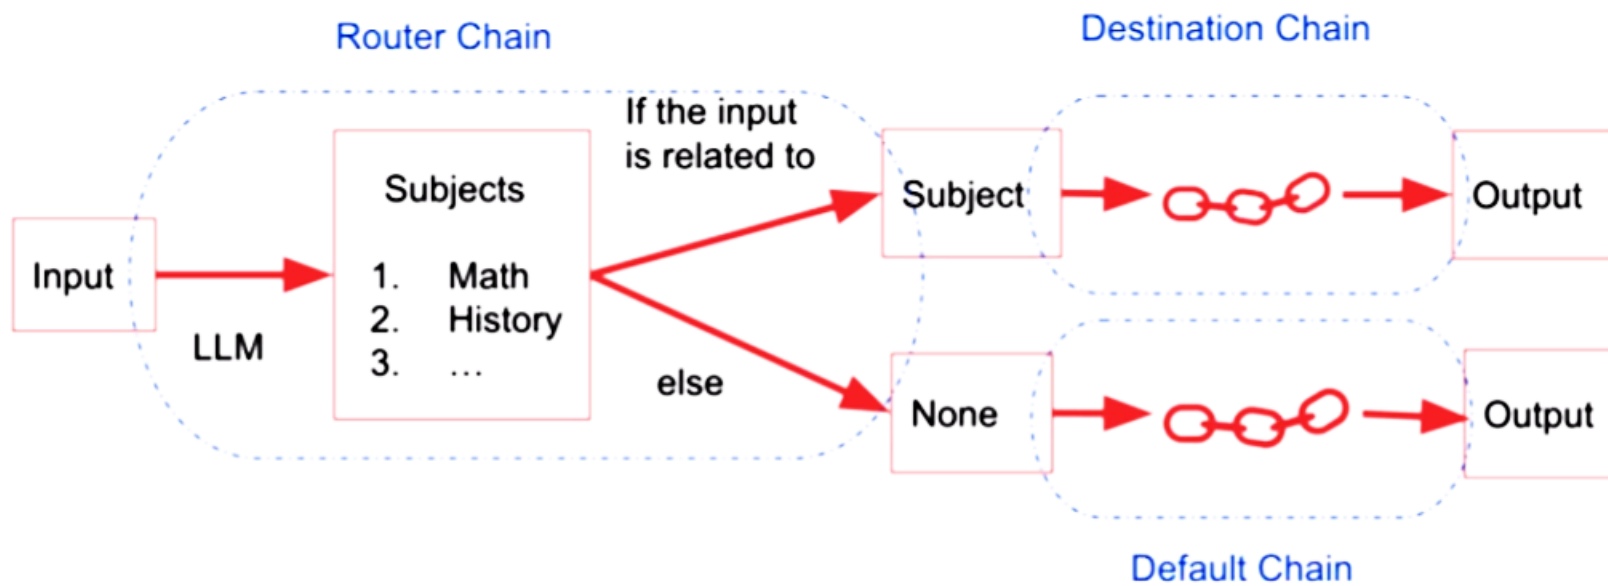


Orchestrating a Solution



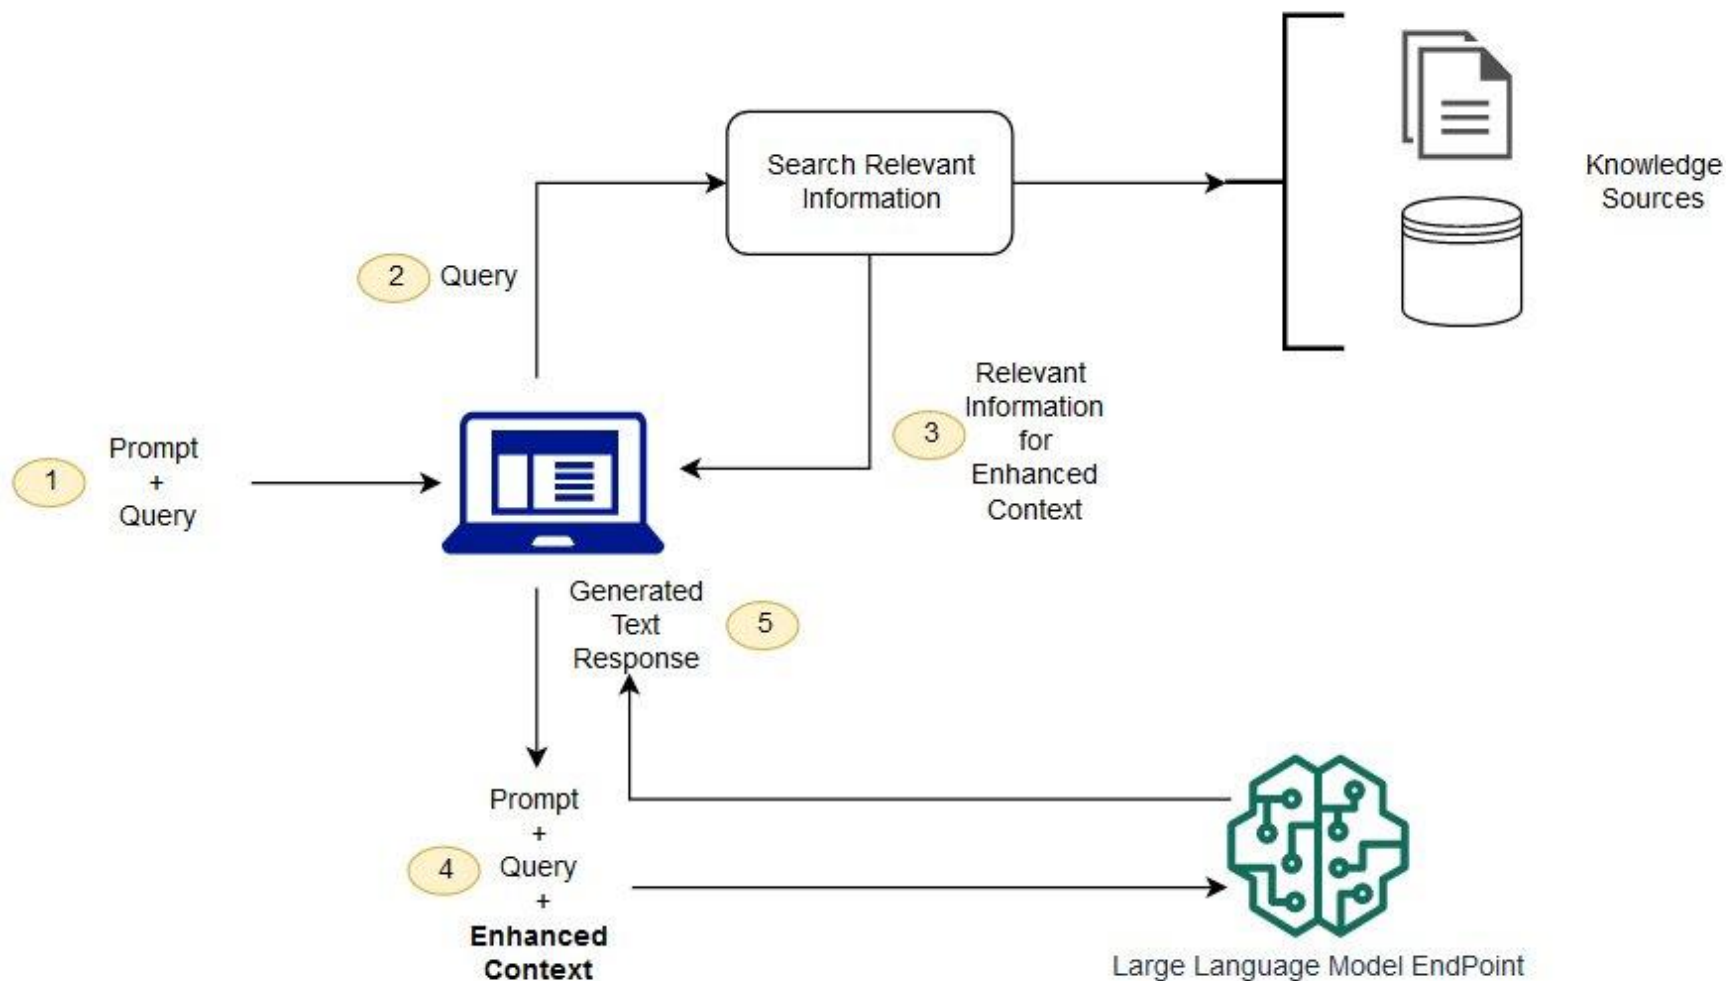


Router Chain



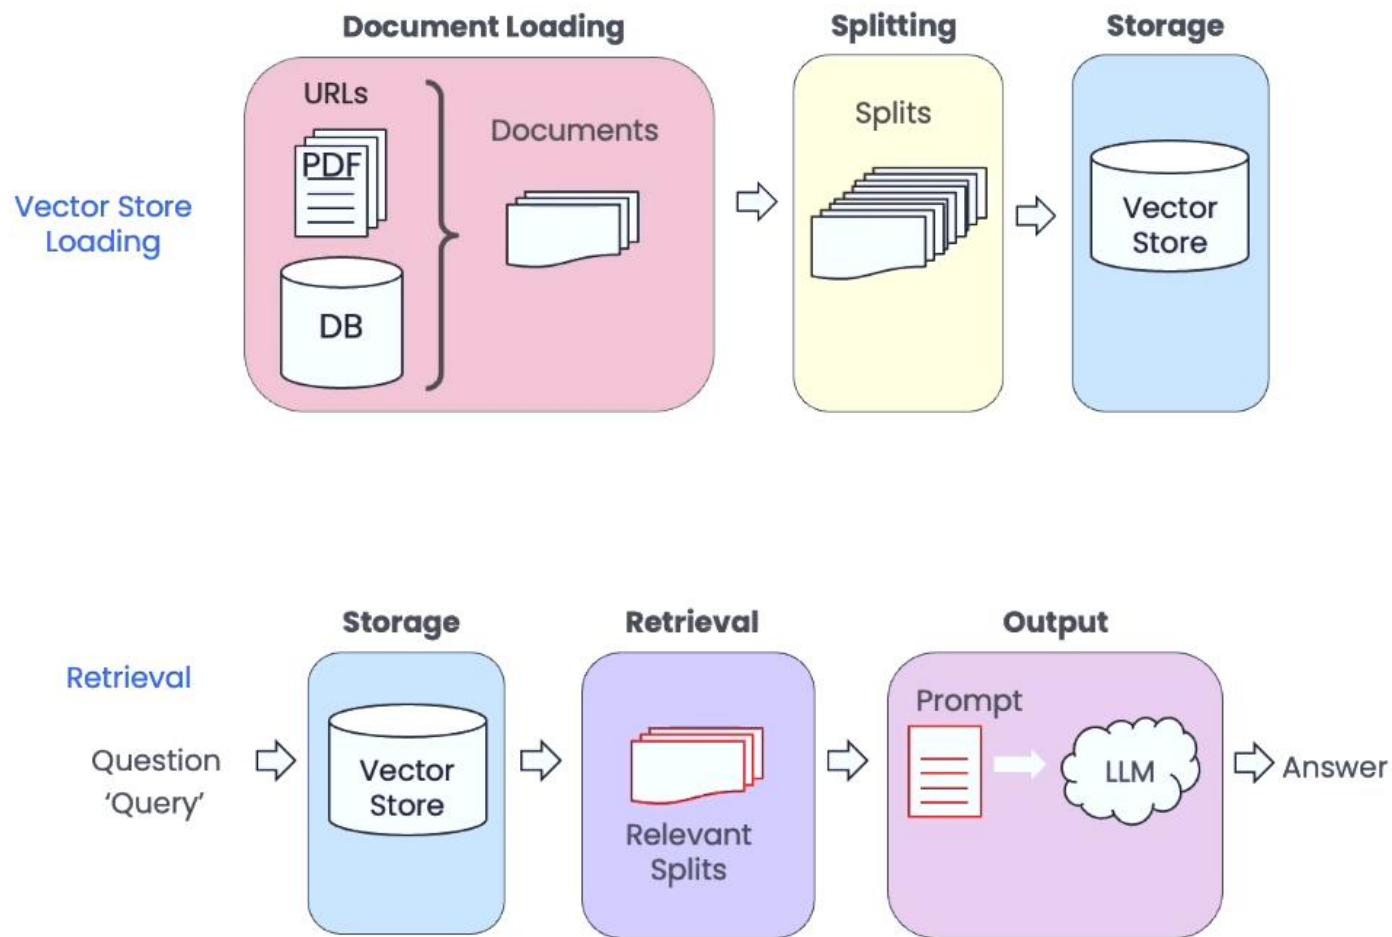


Retrieval Augmented Generation



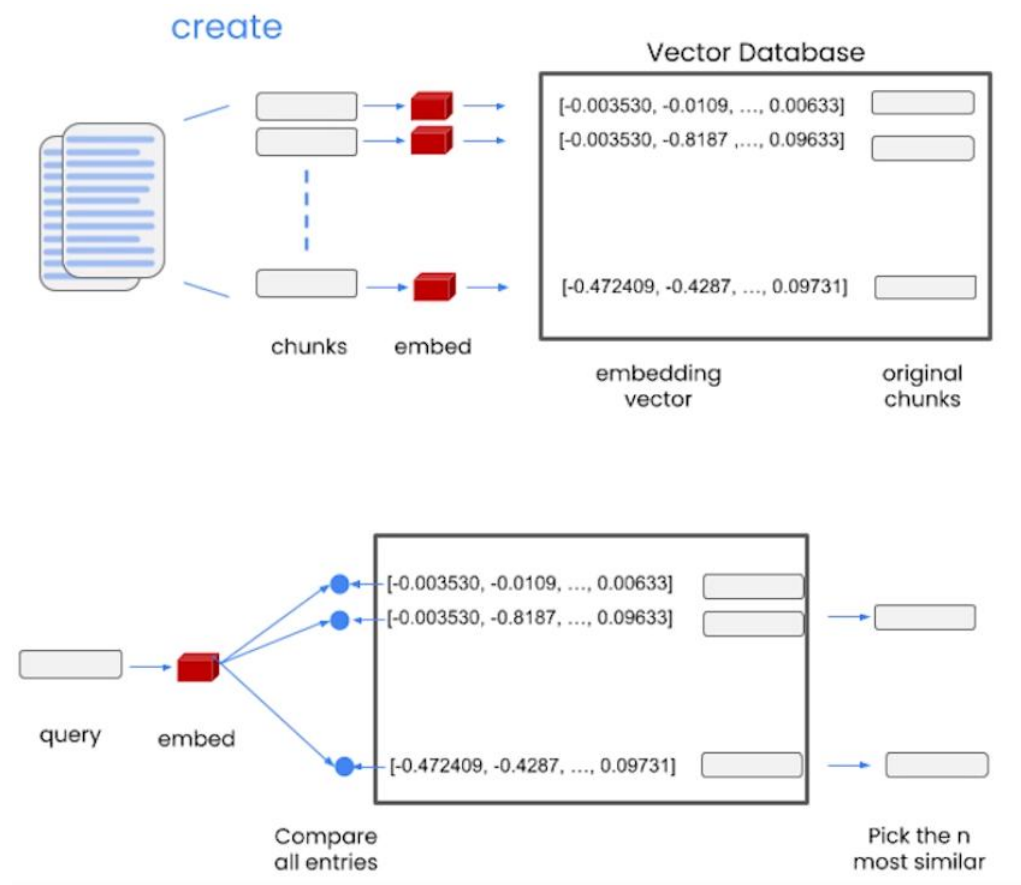


Vector Store and Retrieval





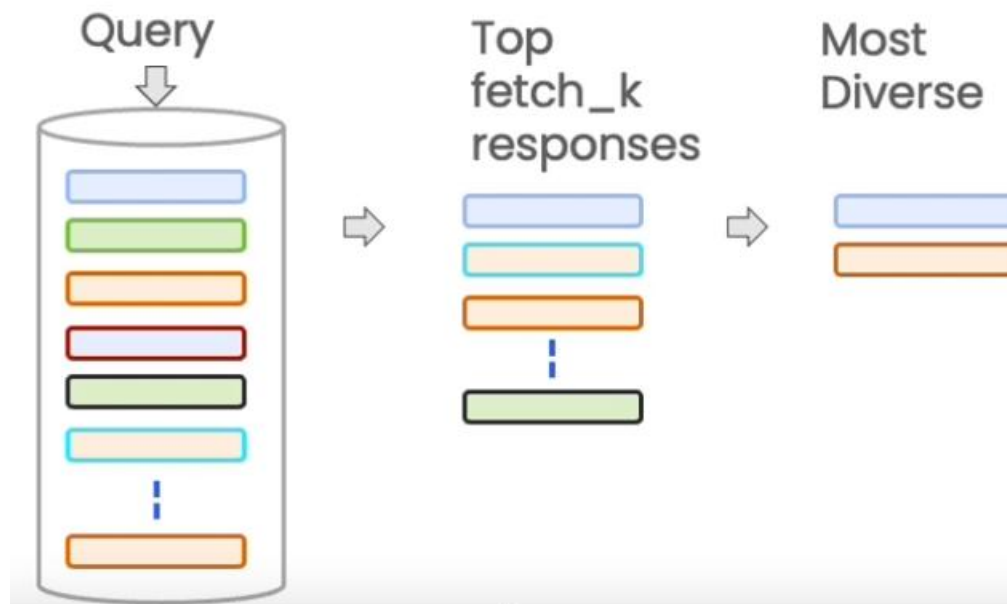
Retrieval in Action



- Data is chopped into chunks and embedded in the vector database
- When a query comes in, the query embedding is compared to all entities and top similar, yet diverse matches are selected



Maximum Marginal Relevance(MMR)



Vector Store Vendors

	Released	Billion-scale vector support	Approximate Nearest Neighbor Algorithm	LangChain Integration
Open-Sourced				
Chroma	2022	No	HNSW	Yes
Milvus	2019	Yes	FAISS, ANNOY, HNSW	
Qdrant	2020	No	HNSW	
Redis	2022	No	HNSW	
Weaviate	2016	No	HNSW	
Vespa	2016	Yes	Modified HNSW	
Not Open-Sourced				
Pinecone	2021	Yes	Proprietary	Yes



Department of Computational and Data Sciences

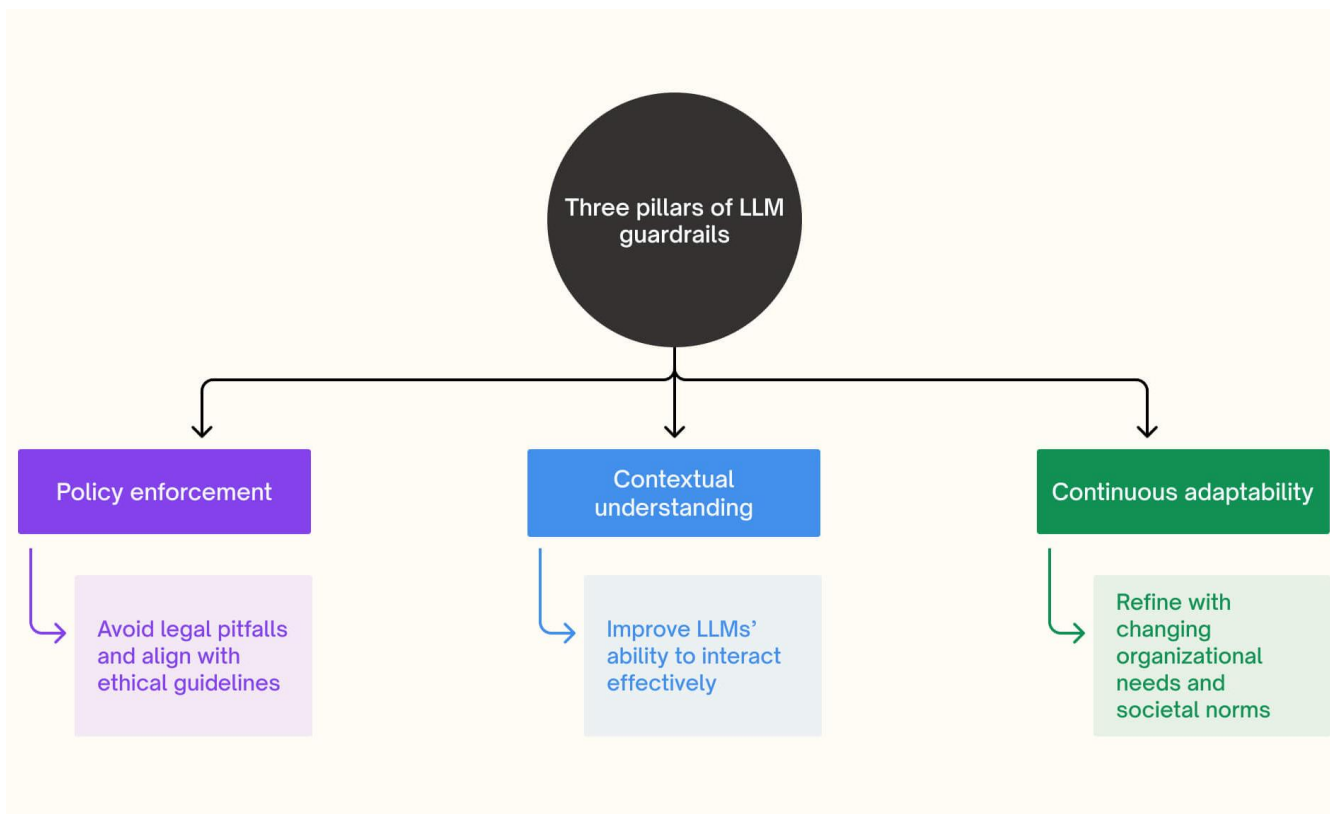
Outline for Feb 01



- Part 1:
 - Instruction Tuning
- Part 2:
 - Orchestration
 - Retrieval Augmented Generation
- Part 3:
 - LLM Guardrails
 - LLM Agents



Why LLM Guardrails are Needed?






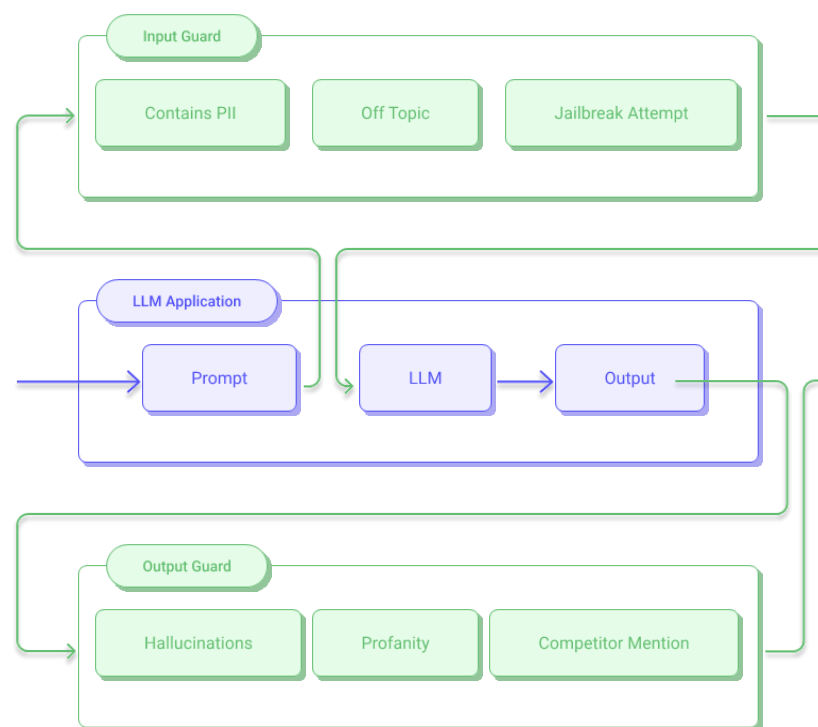
LLM Guardrails



Without Guardrails



With Guardrails



How to Implement Guardrails in your Orchestration?

- Use Guardrails AI
 - <https://hub.guardrailsai.com/>
 - <https://github.com/guardrails-ai/guardrails?tab=readme-ov-file>
- OpenAI's Moderation AI
 - <https://platform.openai.com/docs/guides/moderation>
- NVIDIA's NeMo Guardrails
 - <https://github.com/NVIDIA/NeMo-Guardrails>

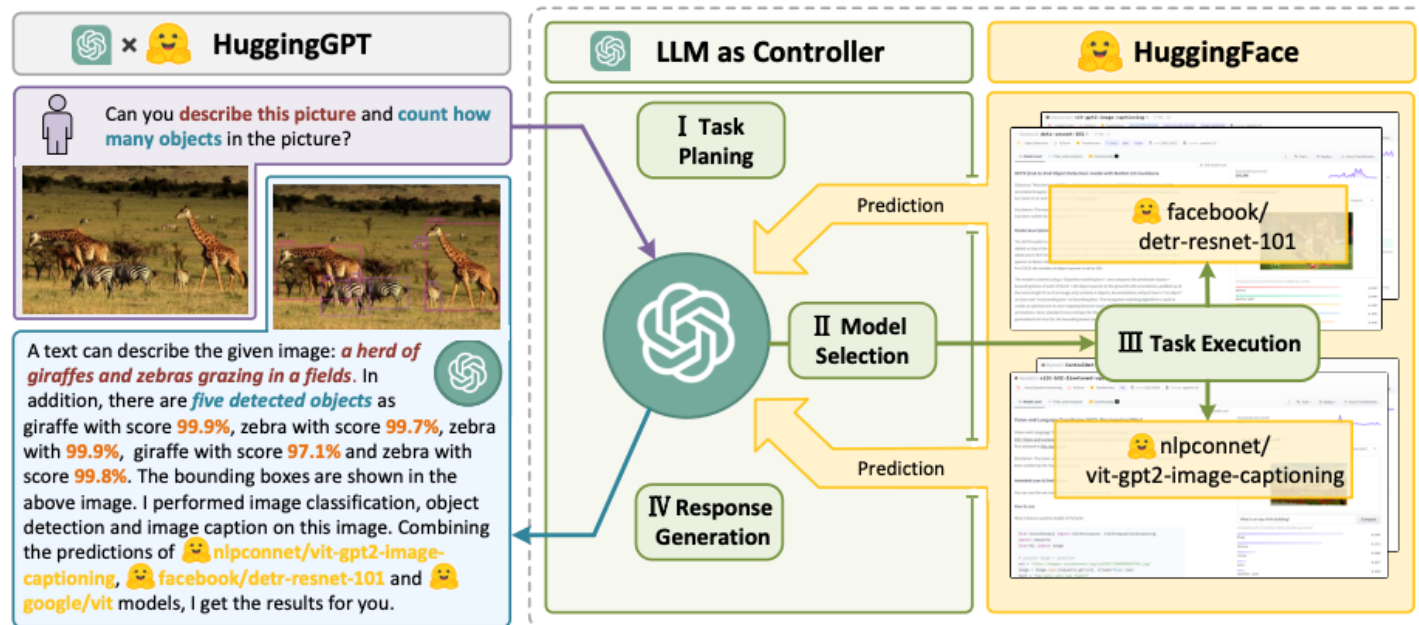
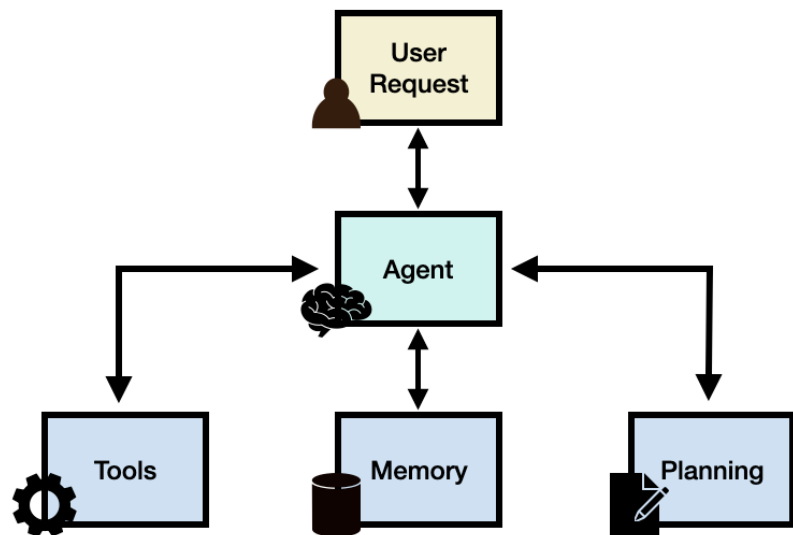


LLM Agents

- LLM Agents are advanced AI systems that use LLMs to understand and generate human language.
- LLM Agents go beyond simple text generation.
 - They can maintain the thread of a conversation,
 - recall previous statements, and
 - adjust their responses accordingly with different tones and styles.
- Diverse Use Cases: customer service, copywriting, data analysis, education, healthcare
- Caution: do not understand nuanced human emotions, and are subject to the risk of misinformation, bias, privacy data leaks and toxicity



LLM Agent Framework





Mixture of Experts Model

- **Gating network** : decides what expert to use

g_1, g_2, \dots, g_k - gating functions

