

DataOps

Data Engineering Fundamentals

Data Engineering Fundamentals

“



Data
Sources
Data



Data
Formats



Data
Models



Storage
Engines &
Processing



Modes of
Dataflow



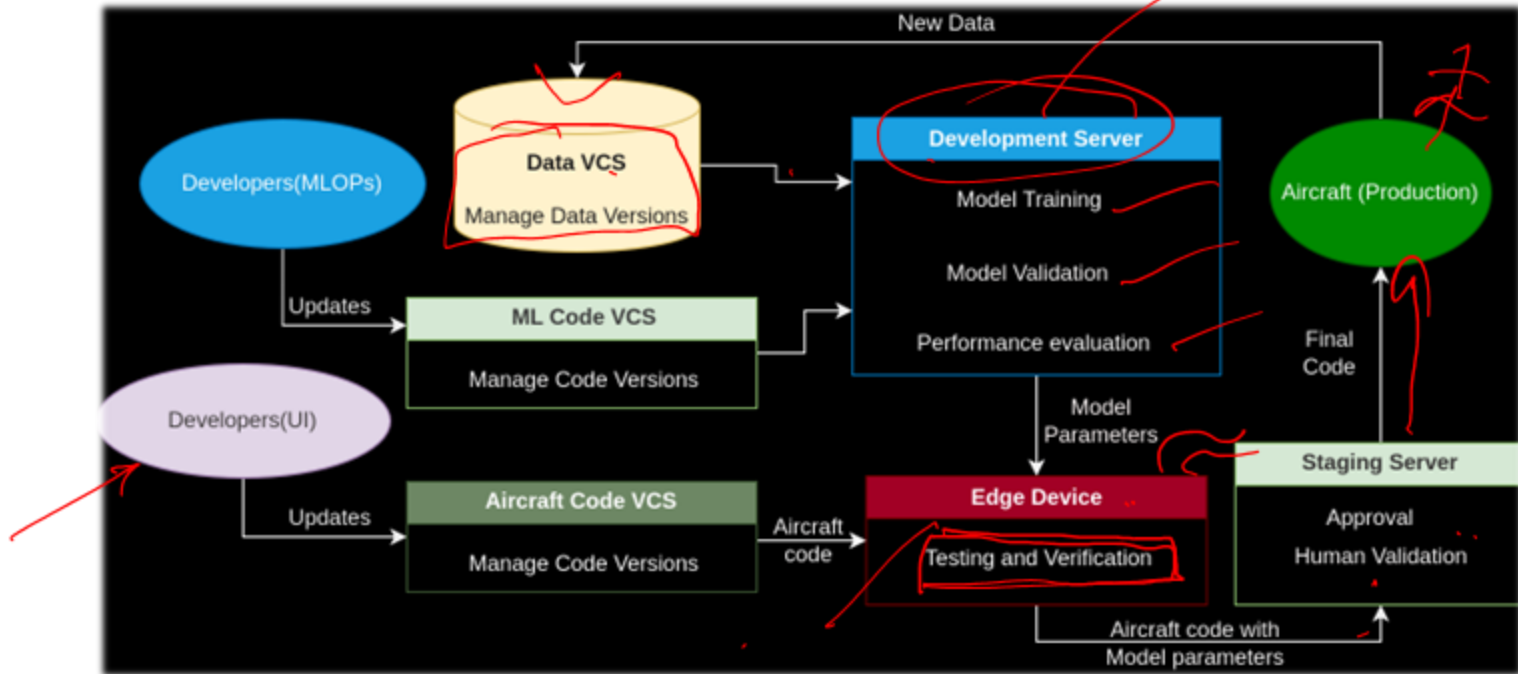
Data Engineering

AI-Agent – Atypical architecture



Data Engineering

AI-Agent – Atypical architecture



“Data Sources

Data Engineering

Data

- something raw, meaningless
- structured, unstructured
- numbers, text, symbols, pixels



Data Engineering

Database

- If data models describe the data in the real world, database specify how the data should be stored on machine
- Storing data is only interesting if we intend to retrieve it later
- Need to know not only how it's formatted but also structured
- Use cases
 - Feature engineering, prediction services, etc.
- Types of Data (storage):
 - Historical data in data storage engines
 - Streaming data in real-time transports

Data Engineering

Data Sources

- ML system needs to work with many different sources with different characteristics, purposes and different processing methods
- User Input:
 - Text, image, video, files, etc
 - High risk for errors
 - User requires fast response, thus needs to be processed quickly
- System-Generated Data
 - Data generated by different components of the system:
 - results, log, system outputs, user activities, etc.
 - Debugging ML system is hard, so log everything



Data Engineering

Data Sources

- Internal Database
 - Data generated by various services and enterprise application within the entity
 - Inventory, customer relationship, users, sales, etc
 - Could be used in different components of ML systems
- First-party Data
 - Data of users or customers
- Second-party Data
 - Collected by another company on their own customers and made it available to us

• Third-party Data

“Data Formats



Data Formats

- Storing data isn't always straightforward since data comes from multiple sources with different access patterns
- Challenges:
 - How to store different type of data, for e.g., multimodal data?
 - Where to economically store and how to access with low latency?
 - How to store complex data so that it can be loaded and run correctly on different hardware?

Data Formats

Format	Binary/Text	Human-readable	Example use cases
JSON	Text	Yes	Everywhere
CSV	Text	Yes	Everywhere
Parquet	Binary	No	Hadoop, Amazon Redshift
Avro	Binary primary	No	Hadoop
Protobuf	Binary primary	No	Google, TensorFlow (TFRecord)
Pickle	Binary	No	Python, PyTorch serialization

Data Formats

JSON

- JavaScript Object Notation
- Most programming languages can generate and parse
- Painful to change the schema
- Human readable
- More space needed since text file

```
{  
  "name": "John Doe",  
  "age": 30,  
  "email": "johndoe@example.com",  
  "address": {  
    "street": "123 Main Street",  
    "city": "New York",  
    "state": "NY",  
    "zip": "10001"  
  },  
  "hobbies": ["reading", "traveling", "photography"],  
  "isStudent": false  
}
```

Data Formats

Row-major Vs. Column-major

- CSV – Comma-separated values

- Text format and row-major
- Good for accessing samples & w
- Low performance

- Parquet – Hadoop, Amazon Reds

- Binary format and column-major
- Highly optimized for big data processing
- Offers efficient compression and encoding schemes
- Good for accessing features and read
- Up to 2X faster to upload and needs 6X less storage in Amazon

S3, compared to text

employee.parquet:

Name	Age	Department
John Doe	30	Engineering
Jane Smith	28	HR
Alex Brown	32	Sales

Data Formats

NumPy Vs. Pandas

- NumPy

- an n-dimensional array object called ndarray, which is a homogeneous, fixed-size array
- Major order can be specified, when created, it's row-major by default

- Pandas

- Two primary data structures:
 - Series and DataFrame.
- Series: A one-dimensional labeled array capable of holding any data type. It is similar to a column in a spreadsheet or a database table.
- DataFrame: A two-dimensional labeled data structure with

Data Formats

Text

(2ms)

- Human-readable: Text data is represented as plain text (e.g., ASCII, UTF-8).
- Interoperability: Text files can be opened and processed by various software applications and programming languages
- Portability: Text files can be easily shared and transferred between different systems.
- Larger file size: Textual representation may require more space to store the same information compared to binary formats.
- Slower processing: Parsing and processing text data can be slower compared to binary data due to the need for data

conversion and interpretation. Common text formats include

Data Formats

Binary

- Compact representation: Binary data is represented in a more compact form, often using numeric values, binary codes, or serialized objects.
- Efficient storage: Binary formats typically require less storage space compared to text formats, especially for large datasets.
- Faster processing: Binary data can be processed faster as it can be directly interpreted by machines without the need for data conversion or parsing.
- Limited human-readability: Binary data is not easily readable or interpretable by humans without specialized tools or knowledge.

“Data Models

Data Models

What is a Data Model?

- Data model describes how data is represented
 - An employees can be described using their qualifications, experience, origin, background, etc.
 - Alternatively, employees can also be described using their job profile, team, performance, impact, etc.
- Representation of data not only affects the way the systems are built, but also the problems it can solve
- Major classification
 - Relational Model and NoSQL model



Data Models

Relational Model

- The relational data model is the foundation of modern database systems
- It organizes data into tables consisting of rows and columns, representing entities and their attributes
- Key Concepts:
 - Tables: Data is stored in tables with predefined columns and rows
 - Rows: Each row represents a unique record or instance of an entity
 - Columns: Columns define the attributes or properties of the entities
 - Primary Key: A unique identifier for each row in a table

Data Models

Relational Model

Row-major:

- Data is stored and retrieved row by row
- Good for accessing samples

Column-major:

- Data is stored and retrieved column by column
- Good for accessing features

	Column 1	Column 2	Column 3
Example 1
Example 2
Example 3



Data Models

Relational Model

- Querying Data:
 - SQL (Structured Query Language) is used to retrieve and manipulate data
 - SELECT statement: Retrieves data from one or more tables
 - JOIN operation: Combines rows from two or more tables based on related columns
 - WHERE clause: Filters data based on specified conditions
- Advantages of Relational Data Model
 - Data consistency: Relationships ensure consistent data across tables
 - Flexibility: Easy to add, modify, or remove data without affecting other parts

Data Models

Relational Model

- Limitations of Relational Data Model:
 - Complex relationships: Many-to-many relationships may require additional tables
 - Eg: students can enroll in multiple classes, and each class can have multiple students
 - Performance impact: Join operations on large tables can impact query performance
 - Vertical scalability: Scaling relational databases vertically (adding more resources to a single server) has limitations

- Real-world Examples:

- Examples of popular relational database management systems

(RDBMS): Oracle, MySQL, SQL Server, PostgreSQL



Data Models

NoSQL

- Not Only SQL is a class of database systems that provide flexible data models beyond the traditional relational model
- Designed to handle large-scale distributed data and non-uniform data structures
- Key Concepts:
 - Schema-less: Do not enforce a fixed schema, allowing for dynamic and flexible data structures
 - Document-oriented: Some NoSQL databases store data in documents (e.g., JSON, XML), representing entities as self-contained documents
 - All information about a record is stored in a document, and easy to retrieve



Data Models

NoSQL

- Graph databases
 - NoSQL databases designed for highly connected data, storing entities as nodes and relationships as edges.
- Scalability
 - Built for horizontal scalability, allowing them to handle large volumes of data and high traffic loads.
 - Distributed architecture: Data is distributed across multiple nodes, enabling parallel processing and fault tolerance.
 - Sharding: Data is partitioned and distributed across multiple servers for efficient storage and retrieval.



Data Models

NoSQL

- Data Model Flexibility:
 - Schema-flexible, allowing for agile development and accommodating evolving data structures
- No need for migrations
 - Changes to the data model can be made on-the-fly without requiring complex schema migrations
- Dynamic attributes
 - Documents or entities can have varying attributes, providing flexibility for different data requirements

Data Models

NoSQL

- Use Cases:
 - Big data analytics: Excel at handling large volumes of diverse data
 - Suitable for real-time applications, such as social networks, IoT, and gaming.
 - Document-oriented NoSQL databases are well-suited for content management systems and CMS-like applications.
- NoSQL Databases:
 - MongoDB: A popular document-oriented NoSQL database.
 - Apache Cassandra: A distributed columnar NoSQL database.
 - Neo4j: A graph database for highly connected data.
 - Redis: A key-value store with in-memory caching capabilities.



Data Models

Structured

- Follows a predefined format and is organized into fixed fields and records.
- Typically stored in relational databases with tables, rows, and columns.
- A clear and rigid schema that defines the data types and relationships.
- Enables straightforward data querying, aggregation, and analysis.
- Excel spreadsheets, SQL databases, CSV files.

Vs.

Unstructured Data


- Lacks a predefined format or structure and doesn't fit into traditional databases.
- Can be in the form of text documents, images, videos, social media posts, etc.
- Doesn't adhere to a rigid schema, making it challenging to define and categorize.
- Analyzing unstructured data requires advanced techniques like natural language processing, image recognition, etc.
- Text documents, social media feeds, multimedia content.



Data Models

Structured Vs. Unstructured Data

- Business requirements change over time, committing to predefined schema can become too restricting
- Unstructured allows more flexible storage options.
- Data warehouse:
 - A repository for storing structured data.
- Data Lakes:
 - A repository for storing unstructured data
 - Used to store raw data before processing

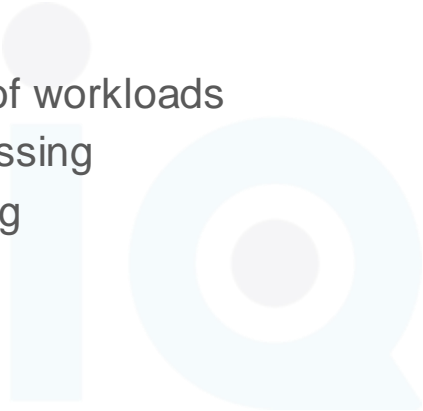


“Data Storage Engines & Processing (Database)



Data Storage Engines & Processing

Storage engine, also known as Database

- An implementation of how data is stored and retrieved on machines
 - Optimized for two types of workloads
 - Transactional processing
 - Analytical processing
- 

Data Storage Engines & Processing

Transactional Processing

- Focuses on managing day-to-day business operations and transactions.
- Involves creating, updating, and deleting individual records in a real-time or near-real-time manner.
- Handle concurrent access from multiple users/applications and ensure consistency.
- ~~ACID~~ properties: Transactional processing emphasizes data integrity and consistency through
 - Atomicity: Ensures all steps are successfully complete, the transaction fails if one step fails
 - Consistency: All transactions must follow predefined rules
 - Isolation: Each transaction is private, no race condition on data access

Data Storage Engines & Processing

Transactional Processing

- Relational model:
 - Transactional systems often use relational databases with well-defined schemas and structured data.
 - OLTP databases: Online Transaction Processing (OLTP) databases are optimized for handling frequent read and write operations.
- Use Cases:
 - Order processing: Managing customer orders, inventory, and payment transactions in e-commerce systems.
 - Banking transactions: Handling account transfers, withdrawals, and deposits in banking systems.
 - Point of sale (POS) systems: Processing sales transactions in



Data Storage Engines & Processing

Transactional Processing

- Focuses on gaining insights and understanding trends from large volumes of data.
- Involves complex queries, aggregations, and data manipulations for decision-making and business intelligence.
- Read-intensive operations: Analytical processing primarily involves querying and analyzing data without frequent updates or modifications.
- Complex queries: Analytical systems execute complex queries involving aggregations, joins, and calculations on large datasets.



Data Storage Engines & Processing

Transactional Processing

- OLAP databases: Online Analytical Processing (OLAP) databases or data warehouses are commonly used for analytical processing.
- Use Cases:
 - Business intelligence: Analyzing sales data, customer behavior, and market trends to make informed business decisions.
 - Data mining: Identifying patterns, correlations, and anomalies in large datasets for predictive analytics and pattern recognition.
 - Reporting and dashboards: Generating summary reports and interactive dashboards to visualize and communicate data insights.

Data Storage Engines & Processing

Analytical Processing

- OLTP and OLAP have become outdated
 - Separation of transactional and analytical databases was due to limitation of technology
 - Was hard to have databases that could handle both queries
 - We have transactional databases that can handle analytical queries and vice versa
 - CrockroachDB
 - Apache Iceberg, DuckDB
 - Storage and processing are tightly coupled
 - May result in data being stored in multiple databases to handle different queries
 - Decouple storage and processing: Google BigQuery, Snowflake

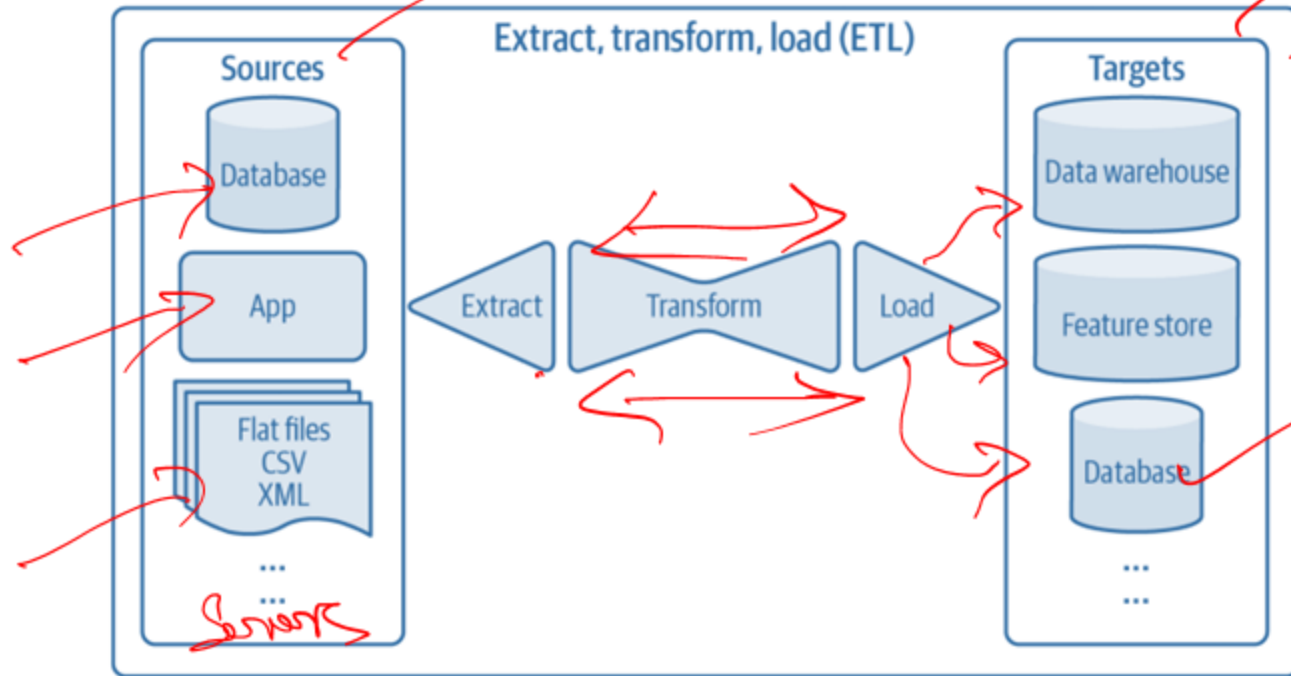
Data Storage Engines & Processing

Extract, Transform, and Load

- A process used in data integration and data warehousing
 - to extract data from various sources
 - transform it into a desired format
 - and load it into a target system for analysis, reporting, or other purposes
- Essential for integrating data from disparate sources, consolidating it into a unified format, and making it available for analysis, reporting, and decision-making.
- Workflows needs to be automated, allow for regular data updates and maintaining data consistency over time, to build ML Systems.

Data Storage Engines & Processing

Extract, Transform, and Load



Data Storage Engines & Processing

Extract

- Data is gathered or retrieved from multiple sources, which can include databases, files, APIs, web services, or external systems.
- Extracted data is typically in its raw or source form, reflecting the structure and format of the source system.
- Data might be corrupt and mal-formatted
- Validation is must, and reject the data that doesn't meet the requirement
- Any wrongdoing affects the entire ETL process

Data Storage Engines & Processing


Transform

- Extracted data is processed and transformed to meet the requirements of the target system or data model.
- Involves cleaning, filtering, aggregating, merging, or modifying the extracted data to ensure consistency, compatibility, and quality.
- May include data cleansing (removing duplicates, handling missing values), data aggregation (summarizing data, calculating metrics), data enrichment (merging with external datasets), and data validation (verifying data integrity, applying business rules).
- Might need to join from multiple sources

Data Storage Engines & Processing

Load

- Transformed data is loaded into the target system, which can be a data warehouse, data mart, analytical database, any other storage system or ML system.
- Loading involves inserting or updating the transformed data in the target system, ensuring data integrity and maintaining relationships.
- The target system is designed for efficient querying, reporting, and analysis of the loaded data.



“Modes of Dataflow



Modes of Dataflow

Main Modes

- Data Passing through Database
 - Data is stored and transferred through relational or non-relational databases.
 - Databases provide persistent storage and structured data management.
 - Data is accessed using SQL or NoSQL queries.
 - Examples: Oracle, MySQL, MongoDB.
- Data Passing through Services
 - Data is exchanged through web services or APIs.
 - Services provide a standardized way of communication between systems.



Modes of Dataflow

Main Modes

- Data Passing through Real-time Data Transport
 - Data is streamed or transferred in real-time for immediate processing or analysis.
 - Real-time data transport enables low-latency data delivery.
 - Data is sent continuously or in small batches.
 - Examples: Messaging systems (Kafka, RabbitMQ), real-time streaming platforms (Apache Kafka, Apache Flink).



Modes of Dataflow

Use Cases

- Database Mode
 - Traditional business applications, transactional systems, data warehousing.
- Service Mode
 - Integrating systems, microservices architecture, API-based interactions.
- Real-time Mode
 - IoT applications, real-time analytics, event-driven systems.
- Hybrid Approaches:
 - Often, multiple modes of dataflow are combined to meet specific requirements.



Modes of Dataflow

To be considered

- Data Volume:
 - Consider the volume of data being transferred and ensure scalability.
- Data Consistency:
 - Maintain data consistency across different modes of dataflow.
- Data Security:
 - Implement security measures to protect data during transfer.
- Latency Requirements:
 - Choose the appropriate mode based on the required data processing speed.



Data Processing

Batch Processing Vs. Stream Processing

- Once data arrives in data storage engines like databases, data lakes, or data warehouse, it become historical data
- Batch Processing:
 - Historical data is often processed in batch jobs.
 - Periodically, for e.g., once a day to compute average surge charge change.
- This is opposite to streaming data since the data is still streaming in.
 - When done right, streaming processing can give low latency (no need to write to database)



Data Processing

Batch Processing Vs. Stream Processing

- Processed in fixed-size batches
 - Collected over a period and processed together as a batch
 - Typically involves processing large volumes of data at once
 - Typically higher latency
 - Overnight data processing, generating reports
- Processed in real-time as continuous streams
 - Processed as it arrives, piece by piece
 - Supports real-time analytics, immediate insights, and near-instantaneous processing
 - Low higher latency
 - Real-time analytics, fraud detection, IoT data

Data Engineering Fundamentals

Summary

- Discussed the importance of data in developing ML Systems
- Learned to choose the right format to store and access data
- Discussed different formats, data models, data engines and processing



Data Sources



Data Formats



Data Models



Data Storage
Engines &
Processing



Modes of
Dataflow

“The Evolving Landscape of Data Storage and Processing

Good to know ...



Data Engineering

What is data engineering?

- The practice of taking raw data from a data source and processing it so it's stored and organized for a downstream use case such as data analytics, business intelligence (BI) or machine learning (ML) model training.
- In other words, it's the process of preparing data so value/insight can be extracted from it.

<https://www.databricks.com/resources/ebook/big-book-of-data-engineering>



Understanding Key Concepts in Modern Data Engineering

Data Warehouse - The Traditional Approach

- A centralized repository for structured data, specifically designed for analytical queries and reporting.
- Key Characteristics:
 - Structured data (relational tables)
 - Pre-defined schema (data is modeled before ingestion)
 - Optimized for read-heavy workloads (BI, reporting)
 - Historical data analysis
- Use Cases:
 - Business intelligence, reporting, dashboards, analytics.



Understanding Key Concepts in Modern Data Engineering

Data Lake - The Rise of Flexibility

- A centralized repository for storing all types of data, both structured and unstructured, in its native format.
- Key Characteristics:
 - Any data type (structured, semi-structured, unstructured)
 - Schema-on-read (data is modeled when queried)
 - Scalable and cost-effective storage
 - Raw data exploration and discovery
- Use Cases:
 - Data science, machine learning, big data analytics, data exploration.
- Challenges:
 - Data governance, data quality, data discovery can be complex.

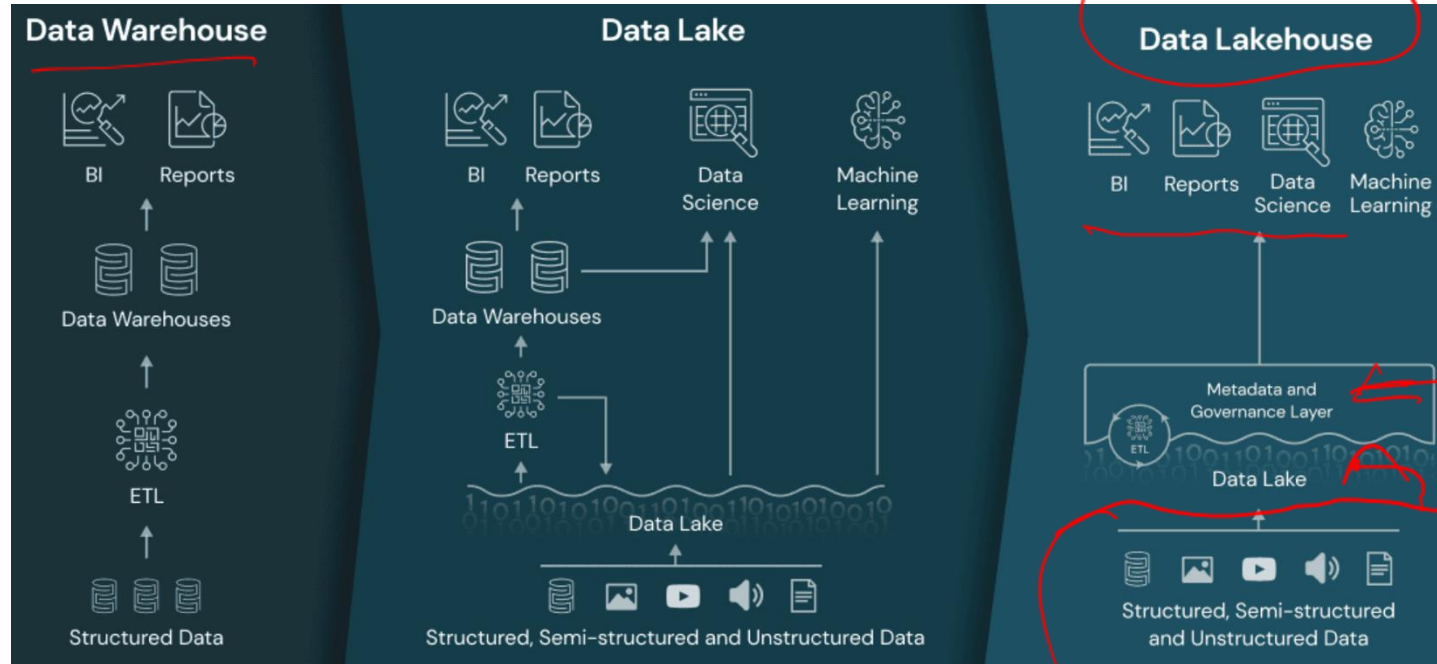


Understanding Key Concepts in Modern Data Engineering

Data Lakehouse - Bridging the Gap

- Combines the scalability and flexibility of a data lake with the data management capabilities of a data warehouse.
- Key Characteristics:
 - Supports both structured and unstructured data
 - Enables ACID (Atomicity, Consistency, Isolation, Durability) transactions and data governance
 - ACID properties are crucial for ensuring data integrity and reliability, especially in applications that handle critical data, such as financial transactions, healthcare records, and e-commerce operations
 - Optimized for both read and write workloads
 - Open and interoperable
- Benefits:
 - Faster time to insights, reduced data duplication, improved data quality.

Understanding Key Concepts in Modern Data Engineering



Understanding Key Concepts in Modern Data Engineering

Delta Lake - The Foundation of the Lakehouse

- An open-source storage layer that brings reliability and performance to data lakes.
- Key Characteristics:
 - ACID transactions for data reliability
 - Schema enforcement and evolution
 - Time travel (data versioning)
 - Unified batch and streaming data processing
 - Optimized for Apache Spark
- Benefits:
 - Improves data quality, simplifies data management, enables real-time

<https://www.databricks.com/wp-content/uploads/2020/08/p975-armbrust.pdf>
https://www.cidrdb.org/cidr2021/papers/cidr2021_paper17.pdf



Understanding Key Concepts in Modern Data Engineering

Data Mesh - Decentralized Data Ownership


- A decentralized approach where different business domains manage their own data as self-service data products.
- Key Features:
 - Domain-oriented design, federated governance, data-as-a-product
- Benefits:
 - Scalability, ownership, agility.
- Example Tools:
 - DataHub, OpenMetadata, Starburst



Understanding Key Concepts in Modern Data Engineering

Data Fabric - Unified Data Access

- An integrated data architecture that enables real-time and secure data access across hybrid and multi-cloud environments.
- Key Features:
 - Unified data access, automation, security.
- Benefits:
 - Real-time insights, improved data governance, simplified data management.
- Example Tools:
 - IBM Data Fabric, Talend, Informatica



Data Engineering

Components and Tools

- **Data Ingestion**

- Data ingestion is the process of bringing data from one or more data sources into a data platform.
- Tools - Apache Kafka, Apache NiFi, Amazon Kinesis.

- **Data Transformation**

- Takes raw ingested data and uses a series of steps (referred to as “transformations”) to filter, standardize, clean and finally aggregate it so it’s stored in a usable way.
- A popular pattern is **the medallion architecture**, which defines three stages in the process — ~~Bronze~~, ~~Silver~~ and Gold.
- Tools - Apache Spark, Dask.

- **Data Storage:** Solutions - HDFS, S3, Google Cloud Storage, Delta Lake.

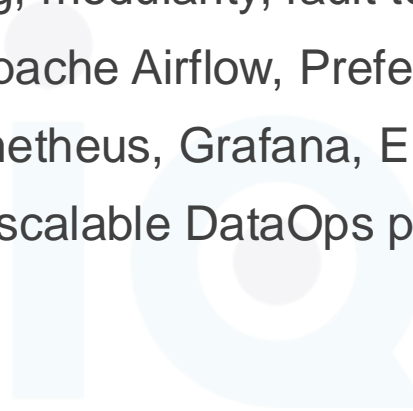
- **Data Quality:** Tools - Great Expectations, Apache Griffin.



Data Engineering

Components and Tools

- **Architecture:** Decoupling, modularity, fault tolerance.
- **Orchestration:** Tools - Apache Airflow, Prefect, Apache NiFi.
- **Monitoring:** Tools - Prometheus, Grafana, ELK stack.
- **Case Study:** Real-world scalable DataOps pipeline example.



Data Engineering

Components and Tools

- **Distributed Training:**

Tools - Horovod, Dask-ML, Spark MLlib.

- **Model Deployment:**

Kubernetes, KFServing, MLflow

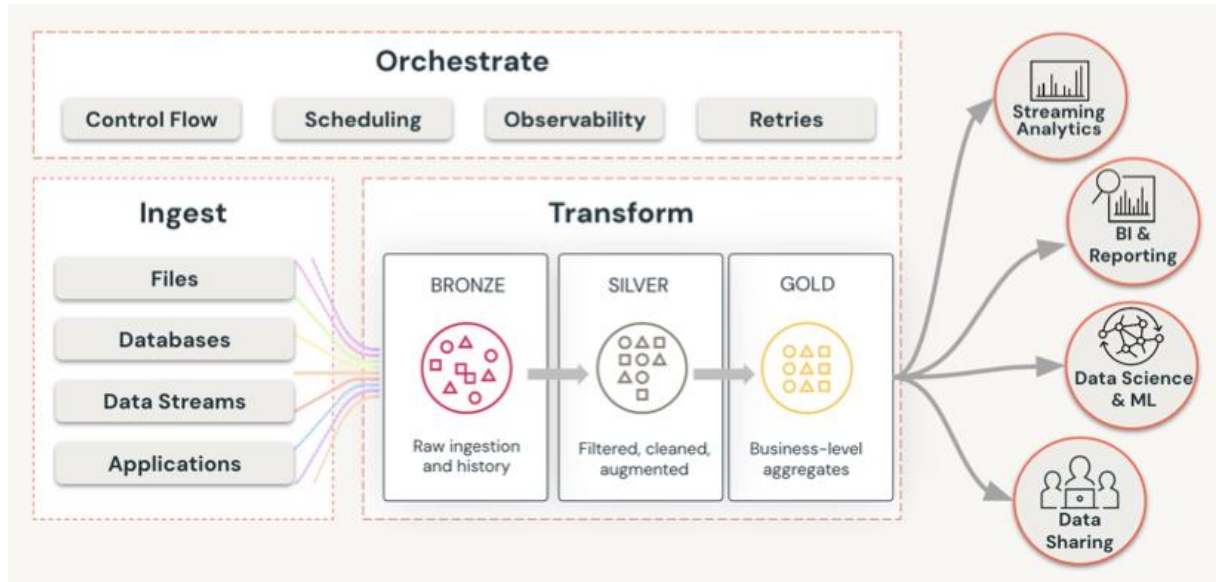
- **Monitoring/Retraining:**

Tools - Evidently AI, Prometheus, Grafana, CI/CD with Jenkins

no Actions

Data Engineering

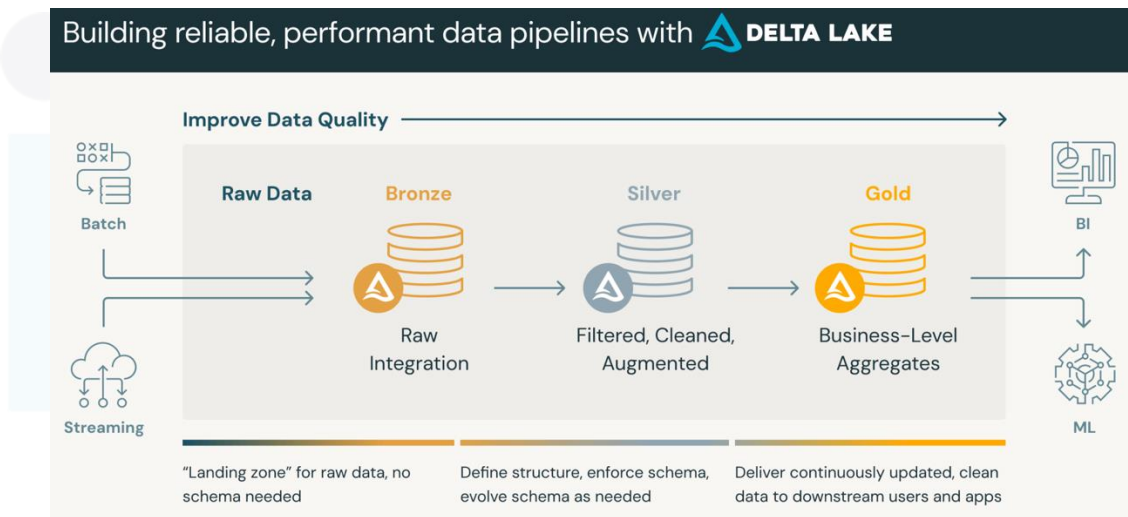
Orchestration



Data Engineering

What is a medallion architecture?

- A data design pattern used to logically organize data in a lakehouse, with the goal of incrementally and progressively improving the structure and quality of data
- Flows through different layers, from Bronze \Rightarrow Silver \Rightarrow Gold layer tables
- Also referred to as "multi-hop" architectures.



Source: databricks



Data Engineering – An Example

Event Streaming

- **Real-Time Processing:** Event streaming includes the manipulation, processing, and real-time reaction to the event streams, enabling immediate responses to data inputs.
- **Retrospective Analysis:** Allows for retrospective processing and analysis of event streams.
- **Data Routing:** Involves routing event streams to different destinations and technologies as required.
- **Continuous Data Flow and Interpretation:** Ensures a consistent flow and interpretation of data, ensuring information is timely and accurately placed.



Data Engineering – An Example

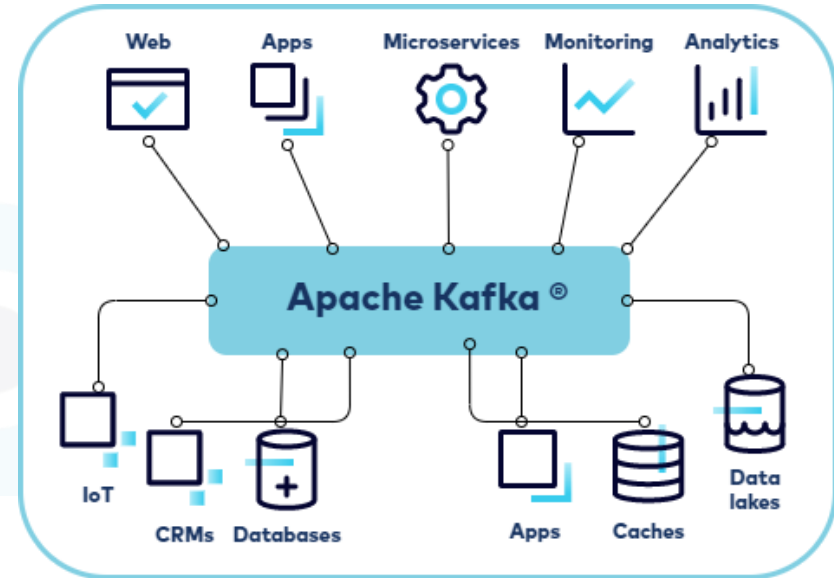
Event Streaming – What for?

- **Financial Transactions:** Used for real-time processing of payments and transactions in sectors like stock exchanges, banking, and insurance.
- **Logistics Tracking:** Applied to track and monitor vehicles in real time, such as cars, trucks, fleets, and shipments, particularly in logistics and the automotive industry.
- **IoT and Sensor Data Analysis:** Utilized to continuously capture and analyze data from IoT devices or equipment in settings like factories and wind parks.
- **Customer Interaction in Retail:** Helps in collecting and responding immediately to customer interactions and orders in retail, hotels, travel industry, and mobile apps.

Data Engineering – An Example

Event Streaming – What for?

- **Healthcare Monitoring:**
Employed in hospitals to monitor patients and predict changes in condition for timely treatment in emergencies.
- **Corporate Data Integration:**
Used to connect, store, and make data available from different divisions within a company.
- **Foundation for Tech Architecture:** Serves as a base for building data platforms, event-driven architectures, and microservices.



Source: <https://docs.confluent.io/kafka/introduction.html>