

# M8-2

## Scalable Deep Learning Algorithms

# Topics



Role of  
Scalability in  
DL's Evolution



Challenges  
and  
Opportunities



Distributed  
training with  
TensorFlow



Summary



Distributed  
training with  
PyTorch

“

# The Role of Scalability in Deep Learning's Evolution



# Role of Scalability in DL's Evolution

- **The Evolution of Deep Learning**

- Deep learning has witnessed exponential growth, driven by advances in computational power and the availability of large datasets.
- This rapid evolution has led to the development of increasingly complex model architectures designed to solve a broader range of problems more accurately.

- **Challenges Posed by Growth**

- **Data Volume:** The explosion in data volume demands algorithms capable of learning from vast datasets without prohibitive increases in training time or resources.
- **Model Complexity:** As models grow in complexity, they require more computational power and memory, posing significant challenges for training and inference.

# Role of Scalability in DL's Evolution

- **The Importance of Scalable Algorithms**
  - Scalable deep learning algorithms are essential for harnessing the full potential of large datasets and complex models.
  - These algorithms enable the efficient utilization of computational resources, ensuring that increases in data volume and model complexity do not lead to intractable increases in computational cost.
- **Objectives of Scalable Algorithms**
  - **Improve Training Speed:** By optimizing computational resource use, scalable algorithms significantly reduce the time required to train complex models on large datasets.
  - **Enhance Model Accuracy:** Scalable algorithms allow for the training of more sophisticated models, leading to improvements in accuracy and performance across various tasks.
  - **Boost Efficiency:** These algorithms ensure that deep learning models can be trained and deployed effectively, even as the scale of data and model complexity grows.

# Role of Scalability in DL's Evolution

## Breakthroughs Enabled by Scalable Algorithms

Scalable algorithms have been pivotal in driving progress across various domains of artificial intelligence, especially in natural language processing (NLP), computer vision, and other areas

### Natural Language Processing (NLP):

- **Efficient Parallel Processing**
  - The Transformer architecture (**Attention Is All You Need**: <https://arxiv.org/abs/1706.03762>, 2017) introduced efficient parallel processing of data, enabling faster training times and improved handling of complex language data compared to traditional sequential models like RNNs and LSTMs.
- **Bidirectional Context Understanding**
  - BERT (<https://arxiv.org/abs/1810.04805>, 2018), built on the Transformer model, advanced NLP by learning from both directions of text simultaneously. This bidirectionality allowed for a nuanced understanding of language context, enhancing performance on tasks requiring deep linguistic comprehension.
- **Leveraging Massive Datasets**
  - Scalability enabled BERT to learn from enormous datasets, including the entirety of

Wikipedia and the BooksCorpus.

# “Distributed training with TensorFlow

# Distributed training with TensorFlow

## MirroredStrategy

### Purpose:

To perform synchronous training across multiple GPUs on a single machine.

### Function:

Each GPU maintains an identical copy of the model. Computes gradients in parallel. **Synchronization:**

Gradients are averaged across all devices before updating the model, ensuring consistency.

### Best used for:

Systems with multiple GPUs seeking to reduce training time without compromising on data integrity.

*[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)*







# Distributed training with TensorFlow

## How MirroredStrategy Works:

### Model Replication:

The strategy replicates the model's variables (weights) across all the GPUs.

### Data Distribution:

The input data is sliced and distributed evenly across the GPUs.

### Gradient Calculation:

Each GPU calculates the gradients for its chunk of the data independently.

### Gradient Aggregation:

The gradients from all GPUs are then gathered and averaged.

### Update Models:

The averaged gradients are applied to the models on all GPUs simultaneously.

*[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)*

# Distributed training with TensorFlow

## How MirroredStrategy Works:

```
import tensorflow as tf

# Assuming you have a model defined as `model`

# Create a MirroredStrategy
strategy = tf.distribute.MirroredStrategy()

# Wrap your model training step in the strategy scope
with strategy.scope():
    # Compile, fit, or train your model using TensorFlow APIs
    model.compile(...)
    model.fit(...) # Or custom training loop
```

*[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)*

# Distributed training with TensorFlow

## TPUStrategy

- **Purpose:**
  - Optimizes training processes for TensorFlow computations on Google's Tensor Processing Units (TPUs).
- **Function:**
  - Abstracts the complexity of distributed training on TPUs, handles data partitioning to leverage TPU cores.
- **Synchronization:**
  - Coordinates gradient updates across TPU cores effectively.
- **Best used for:**
  - Projects requiring extreme computation power and rapid execution, such as large-scale neural networks.

*[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)*

# Distributed training with TensorFlow

## How TPUStrategy Works:

- **Model Placement**
  - Unlike MirroredStrategy that replicates on GPUs, TPUStrategy places the entire model (or a partitioned version) onto the TPU device.
- **Data Parallelism**
  - The strategy splits the input data across the TPU cores.
  - Each core processes a portion of the data.
- **Gradient Aggregation**
  - All TPU cores calculate gradients for their data shard and then the gradients are synchronized across all cores.
  - This means that all cores update their local model weights in a coordinated fashion after each global batch is processed.
- **Performance Optimization**
  - TPUStrategy includes optimizations specific to TPUs, such as utilizing TPU-specific operations for faster execution.

[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)

# Distributed training with TensorFlow

## How TPUStrategy Works:

```
import tensorflow as tf

# Assuming you have a model defined as `model`

# Create a TPUStrategy
resolver = tf.distribute.cluster_resolver.TPUClusterResolver() # Identify TPUs
strategy = tf.distribute.TPUStrategy(resolver)

# Wrap your model training step in the strategy scope
with strategy.scope():
    # Compile, fit, or train your model using TensorFlow APIs
    model.compile(...)
    model.fit(...) # Or custom training loop
```

*[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)*

# Distributed training with TensorFlow

## MultiWorkerMirroredStrategy: Scaling Training Across Multiple Machines

- **Purpose:**
  - Extends MirroredStrategy for synchronous training across multiple workers, each possibly with multiple GPUs.
- **Function:**
  - Similar to MirroredStrategy but operates over the network.
- **Synchronization:**
  - Uses collective communication methods like All-Reduce to update model copies.


### Best used for:

- Large-scale distributed training tasks over multiple machines.





# Distributed training with TensorFlow



## How MultiWorkerMirroredStrategy Works:

- **Cluster Setup:**
  - You have a cluster of 2 machines (nodes), each node potentially equipped with multiple GPUs.
- **MirroredStrategy on Each Worker:**
  - MultiWorkerMirroredStrategy creates a separate MirroredStrategy instance on each node.
  - This means each node replicates the model across its own available GPUs. 1 - 600
- **Data Sharding:**
  - The training dataset is divided (sharded) into smaller chunks and distributed to each node. 501 - 10000
  - This ensures parallelism across nodes.
- **Synchronized Training:**
  - Each node's MirroredStrategy processes its assigned data shard independently.
  - Gradients are calculated locally on each node.
- **Gradient Exchange:**
  - After processing, gradients are exchanged and aggregated efficiently between all nodes using a technique called "all-reduce."

- **Model Update:**

[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)

# Distributed training with TensorFlow

## How MultiWorkerMirroredStrategy Works:

```
import tensorflow as tf

# Assuming you have a model defined as `model`

# Define cluster configuration (replace with actual cluster details)
cluster_resolver = tf.distribute.cluster_resolver.SlurmClusterResolver(...)

# Create a MultiWorkerMirroredStrategy
strategy = tf.distribute.MultiWorkerMirroredStrategy(cluster_resolver)

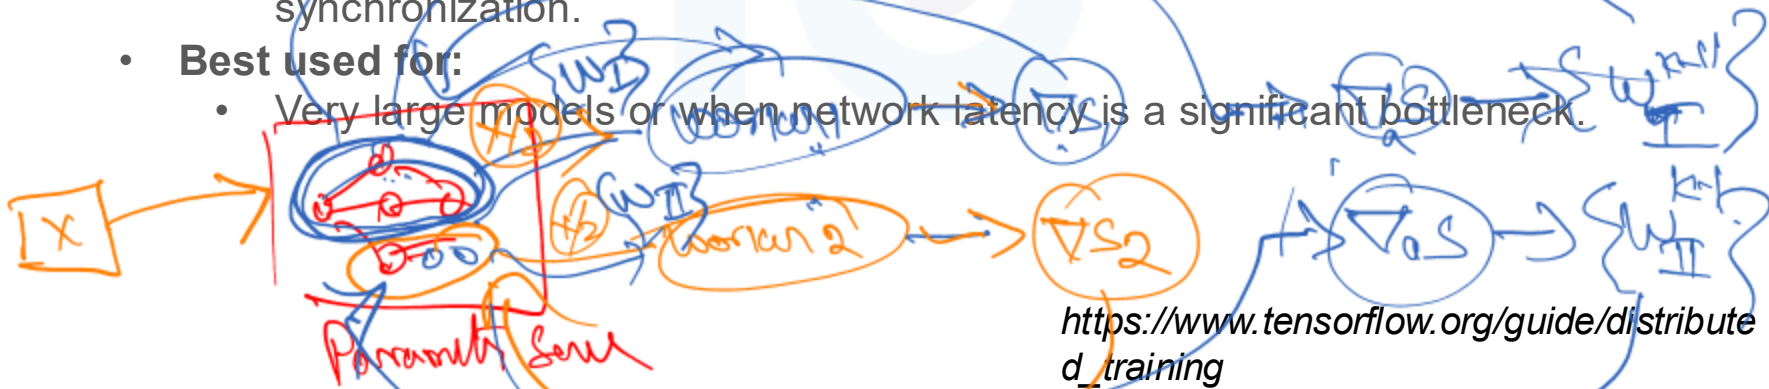
# Wrap your model training step in the strategy scope
with strategy.scope():
    # Compile, fit, or train your model using TensorFlow APIs
    model.compile(...)
    model.fit(...) # Or custom training loop
```

[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)

# Distributed training with TensorFlow

## ParameterServerStrategy (experimental strategy)

- **Purpose:**
  - For asynchronous training with parameter servers.
- **Function:**
  - Model parameters are stored on parameter servers, workers update these asynchronously.
- **Synchronization:**
  - Minimizes network bottleneck, as workers do not need to wait for synchronization.
- **Best used for:**
  - Very large models or when network latency is a significant bottleneck.



[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)

# Distributed training with TensorFlow

## How ParameterServerStrategy Works:

- **Cluster Roles:** The training cluster consists of three types of nodes:
  - **Worker Nodes:**
    - These nodes handle the actual computations for training the model.
    - Each worker can have CPUs or GPUs.
  - **Parameter Servers:**
    - These dedicated servers store and manage the model variables (weights and biases).
    - They don't perform computations directly.
  - **Coordinator: (Optional)**
    - A central coordinator node can be used to manage cluster resources and dispatch tasks to workers.
- **Variable Distribution:**
  - Model variables are sharded (divided) and distributed across multiple parameter servers.
  - This allows for efficient storage and handling of large models.

# Distributed training with TensorFlow

## How ParameterServerStrategy Works:

- **Worker-Server Interaction:**

- During training, each worker node:
  - Downloads a copy of the relevant variable shards from the parameter servers.
  - Performs computations on its assigned data batch using the downloaded variables.
  - Calculates gradients based on the computation results.
  - Sends the calculated gradients back to the parameter servers.

- **Parameter Update:**

- Parameter servers receive gradients from all workers, aggregate them (e.g., averaging), and update their corresponding variable shards.
- This aggregation step is typically **synchronous**. This ensures that all workers contribute to the same update, and the model state remains consistent across the cluster.

- **Potential Alternatives for Asynchronous Parameter Updates:**

- Stale gradients:
  - Workers might use slightly outdated parameter values when calculating gradients.
  - This can improve efficiency but might affect convergence or introduce

# Distributed training with TensorFlow

## How ParameterServerStrategy Works:

```
import tensorflow as tf

# Assuming you have a model defined as `model`

# Define cluster configuration (replace with actual cluster details)
cluster_resolver = tf.distribute.cluster_resolver.SlurmClusterResolver(...)

# Create a ParameterServerStrategy
strategy = tf.distribute.experimental.ParameterServerStrategy(cluster_resolver)

# Wrap your model training step in the strategy scope
with strategy.scope():
    # Compile, fit, or train your model using TensorFlow APIs
    model.compile(...)
    model.fit(...) # Or custom training loop
```

*[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)*

# Distributed training with TensorFlow

## CentralStorageStrategy (experimental strategy)

- **Purpose:**
  - To allow computation to be offloaded from one device to others, while centralizing the model's parameters.
- **Function:**
  - One device holds the parameters, and others assist in computation.
- **Synchronization:**
  - Occurs on the central device, reducing the need for complex communication protocols.
- **Best used for:**
  - Moderate-sized models where device-to-device communication is efficient..

*[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)*

# Distributed training with TensorFlow

## How CentralStorageStrategy Works:

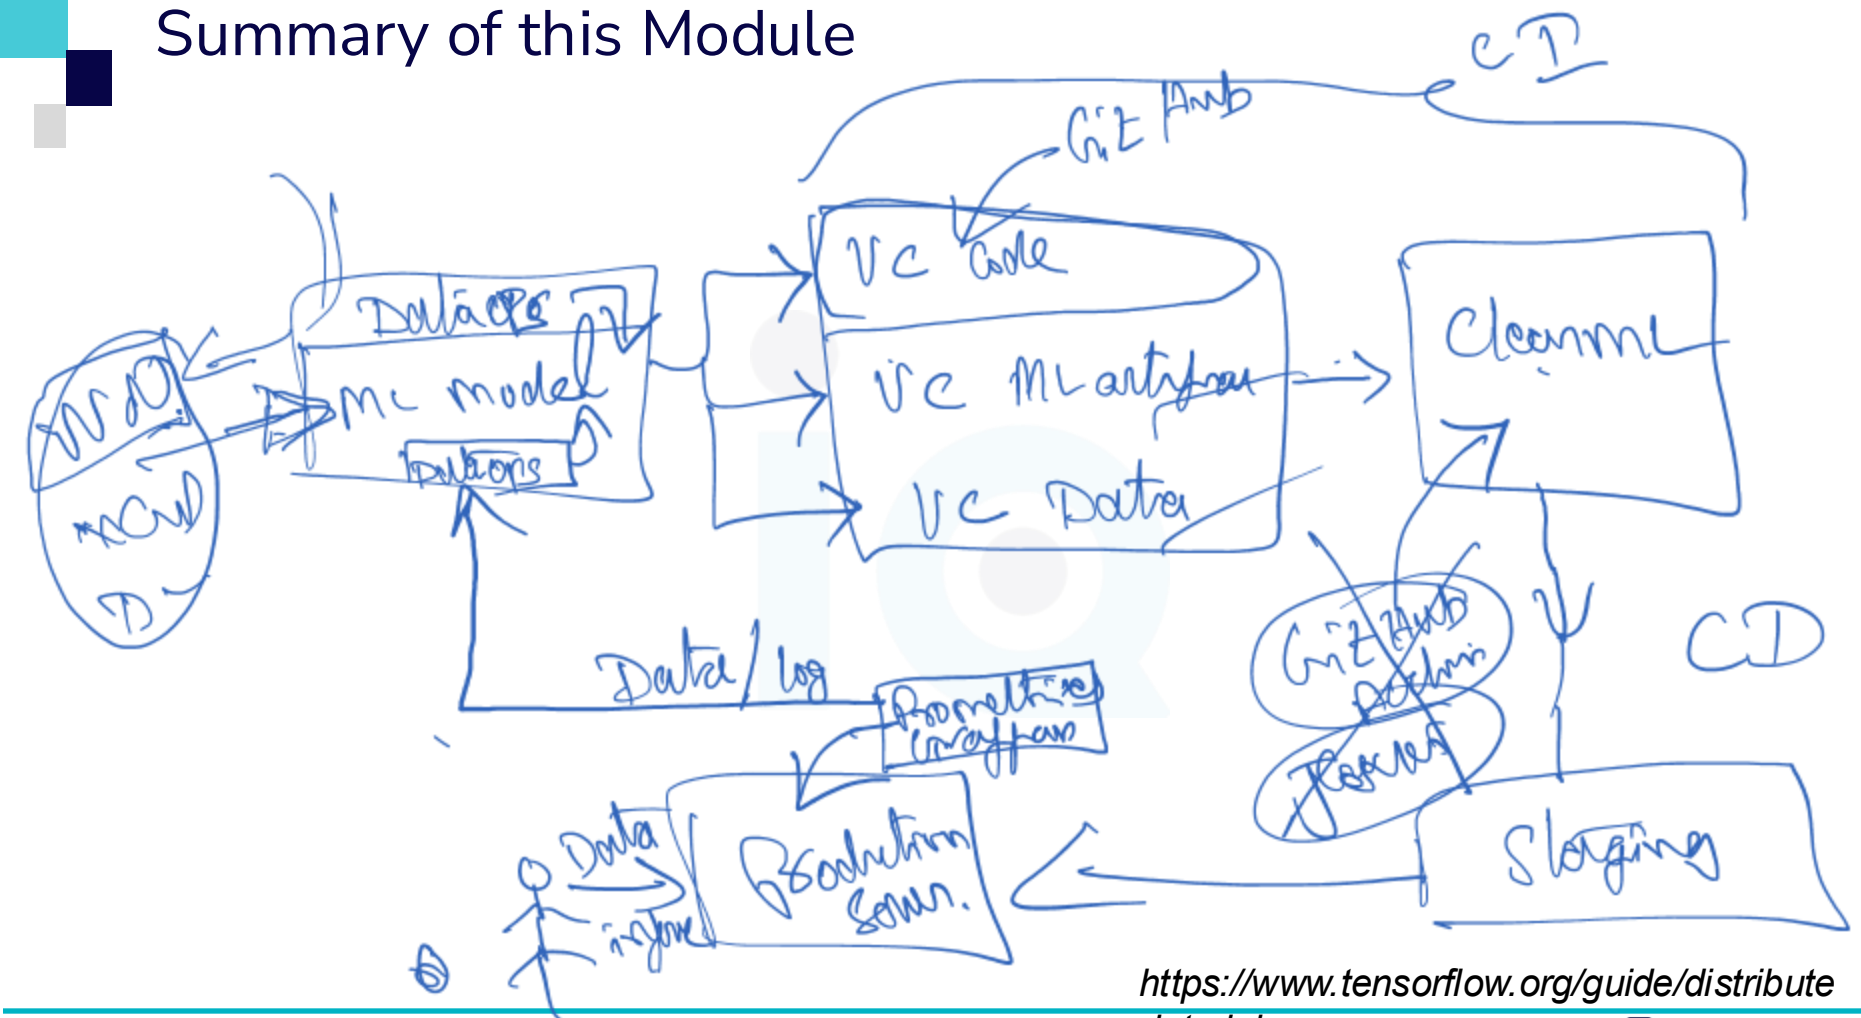
- **Central Storage:**
  - Similar to ParameterServerStrategy, model variables are likely stored centrally on a designated node (CPU or GPU).
- **Worker Computations:**
  - Worker nodes perform computations on their assigned data batches, presumably downloading or accessing the variables from the central storage as needed.
- **Gradient Communication:**
  - Workers calculate gradients and send them back to the central storage for aggregation and model updates.



# Distributed training with TensorFlow

Feature	MultiWorkerMirroredStrategy	ParameterServerStrategy (PS)	CentralStorageStrategy (CS)
<b>Model Replication</b>	Each worker replicates model (uses MirroredStrategy)	Variables distributed across parameter servers	Variables stored centrally (potentially)
<b>Worker Roles</b>	All workers perform computations	Workers compute, parameter servers store variables	Workers compute, central storage for variables
<b>Communication Pattern</b>	All-reduce gradients between workers	Workers communicate with parameter servers	Workers communicate with central storage
<b>Scalability</b>	Good for large datasets, decent for large models	Potentially better for very large models	Potential for scalability, details unclear
<b>Flexibility</b>	Workers and GPUs can be scaled together	Workers and parameter servers can be scaled independently	Limited information on flexibility
<b>Complexity</b>	Relatively simpler to set up	More complex setup and configuration	Experimental, complexity unclear
<b>Synchronous/Asynchronous</b>	Synchronous updates	Typically synchronous parameter updates, potential for asynchronous exploration	Information limited, might be synchronous
<b>Fault Tolerance</b>	Limited fault tolerance	Potentially better due to separate storage	Information limited
<b>Status</b>	Stable	Experimental	Experimental

# Summary of this Module



[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)



# Summary of this Module



[https://www.tensorflow.org/guide/distributed\\_training](https://www.tensorflow.org/guide/distributed_training)