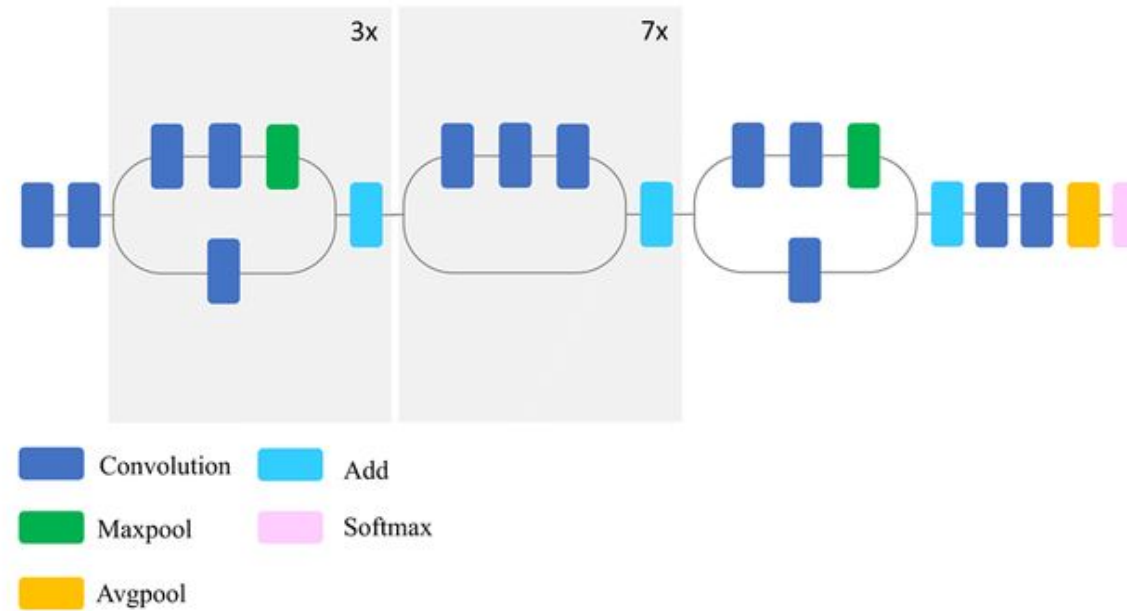


Modern Convnet Architecture



OBJECTIVES

□ Understanding the important building blocks of modern CNNs

- ❖ Residual connections
- ❖ Batch normalisation
- ❖ Depthwise separable convolution

□ Building a Mini-Xception Model

□ Interpret what CNNs learn

1. Visualising activations
2. Visualising filters
3. Visualising heatmaps


For

- 1) Use : Small Convnet from last AST: Cat/Dog classification with Augmentation
- 2) & 3) Use: Pretrained Xception model

PART : A

Understanding the important building blocks of modern CNNs

VGG16 Architecture




- Note the repeated layer blocks and the pyramid-like structure of the feature maps.


- Deeper hierarchies are intrinsically good because they encourage feature reuse, and therefore abstraction. In general, a deep stack of narrow layers performs better than a shallow stack of large layers. However, there's a limit to how deep you can stack layers, due to the problem of vanishing gradients.

□ **Residual connection**

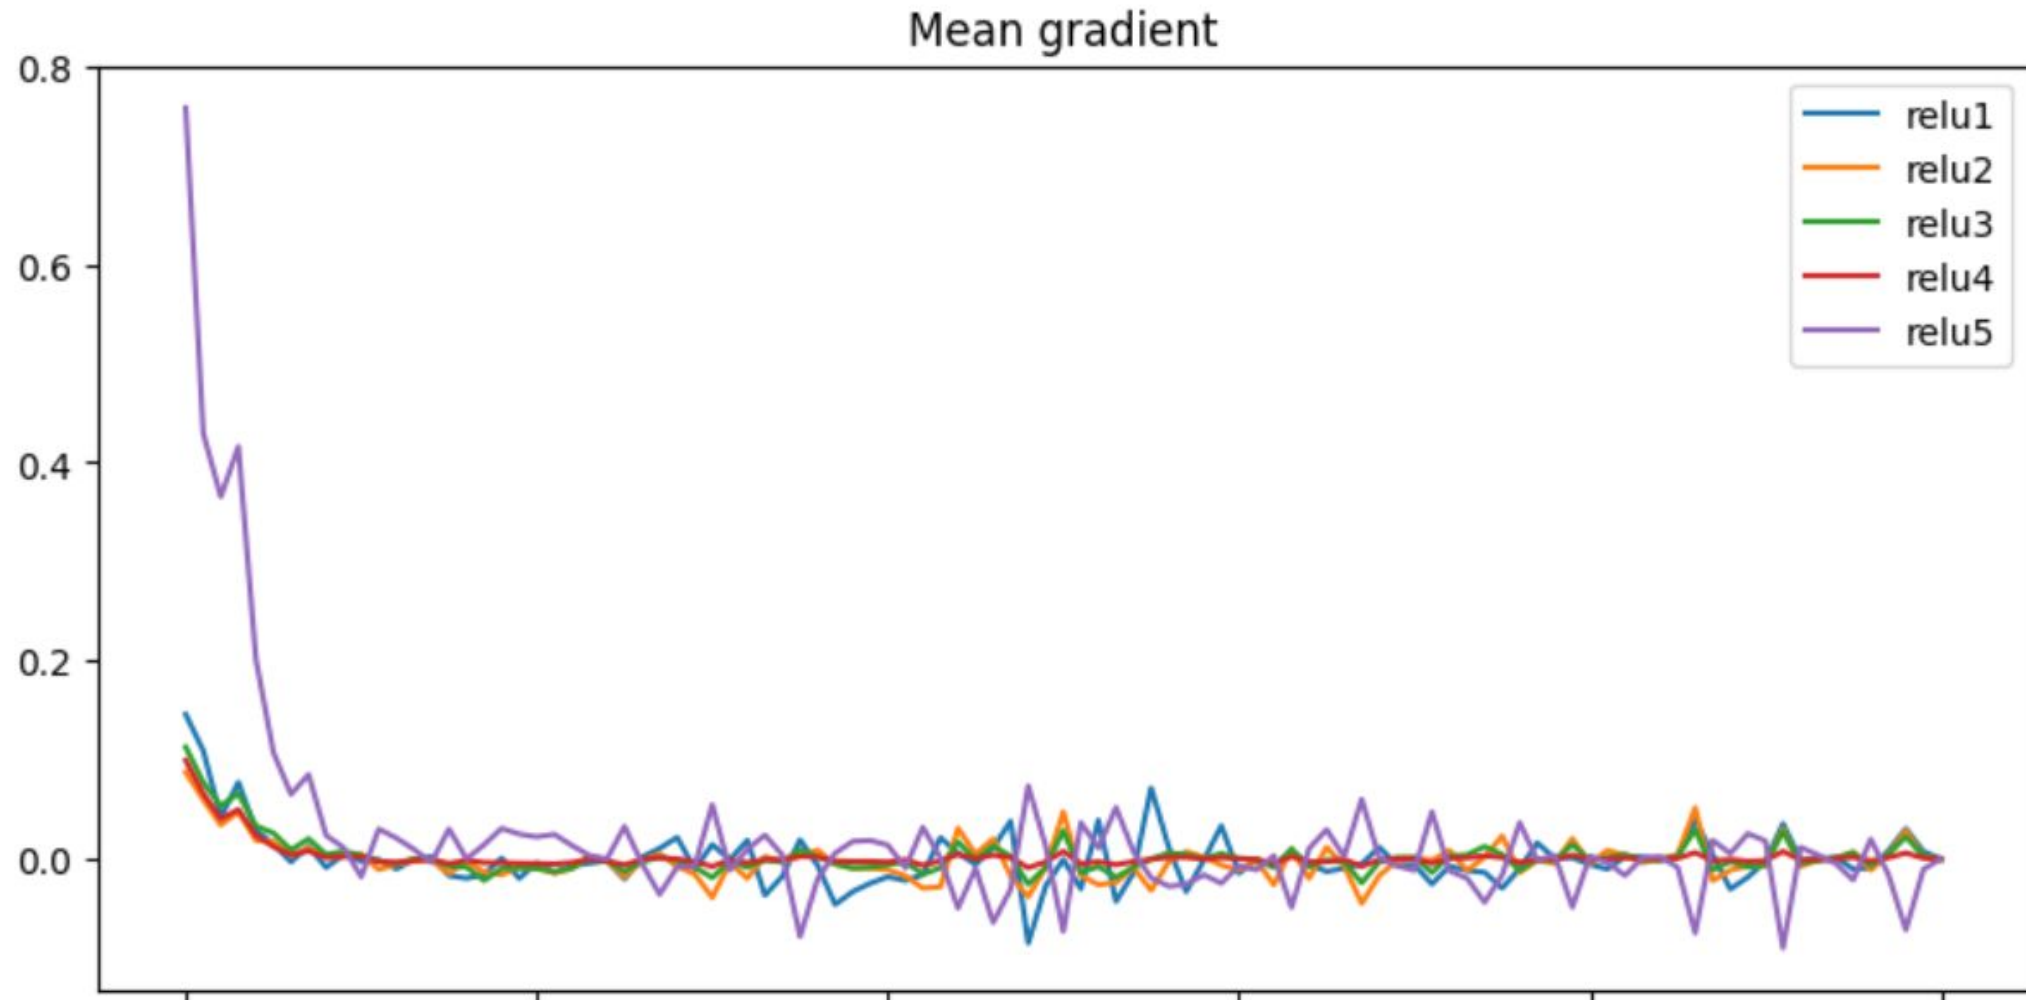
Vanishing gradients - explained



Vanishing gradients



Vanishing gradients



Vanishing gradients

Input gradients are calculated from output gradients using the formula

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial y} \cdot f'(x)$$

where:

$$\left(\frac{\partial L}{\partial x} \right)$$


is the gradient of the loss with respect to the input (x),

$$\left(\frac{\partial L}{\partial y} \right)$$

is the gradient of the loss with respect to the output (y),

$$f'(x)$$


is the derivative of the activation function with respect to x .



Vanishing gradients - solutions

- Weight initialization
- Residual connections
- Batch normalization

A residual connection



A residual connection in pseudocode

```

x = ...
residual = x
x = block(x)
x = add([x, residual])

```

Some input tensor (points to `x = ...`)

Save a pointer to the original input. This is called the residual. (points to `residual = x`)

This computation block can potentially be destructive or noisy, and that's fine. (points to `x = block(x)`)

Add the original input to the layer's output: the final output will thus always preserve full information about the original input. (points to `x = add([x, residual])`)

Case 1: Residual block where the number of filters changes

```
# Build model with residual connection - layer in the residual branch
inputs = keras.Input(shape=(32, 32, 3), name="input")
x = layers.Conv2D(32, 3, activation="relu", name="C1")(inputs)
residual = x
x = layers.Conv2D(64, 3, activation="relu", padding="same", name="C2l1")(x)
x = layers.Conv2D(64, 3, activation="relu", padding="same", name="C2l2")(x)
residual = layers.Conv2D(64, 1, name="C2b")(residual)
x = layers.add([x, residual])
model2 = keras.Model(inputs=inputs, outputs=x)
```

Layer (type)	Output Shape	Param #
input (InputLayer)	[(None, 32, 32, 3)]	0
C1 (Conv2D)	(None, 30, 30, 32)	896
C2l1 (Conv2D)	(None, 30, 30, 64)	18496
C2l2 (Conv2D)	(None, 30, 30, 64)	36928
C2b (Conv2D)	(None, 30, 30, 64)	2112
add_2 (Add)	(None, 30, 30, 64)	0

Case 2: Where the target block includes a max pooling layer

```
# model with residual connections and a max pool layer in between
inputs = keras.Input(shape=(32, 32, 3))
x = layers.Conv2D(32, 3, activation="relu")(inputs)
residual = x
x = layers.Conv2D(64, 3, activation="relu", padding="same")(x)
x = layers.MaxPooling2D(2, padding="same")(x)
residual = layers.Conv2D(64, 1, strides=2)(residual)
x = layers.add([x, residual])
model3 = keras.Model(inputs=inputs, outputs=x)
plot_model(model3, show_shapes=True)
```

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 32, 32, 3)]	0
conv2d_5 (Conv2D)	(None, 30, 30, 32)	896
conv2d_6 (Conv2D)	(None, 30, 30, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 15, 15, 64)	0
conv2d_7 (Conv2D)	(None, 15, 15, 64)	2112
add_4 (Add)	(None, 15, 15, 64)	0

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

The BatchNormalization layer can be used after any layer—Dense, Conv2D, etc.:

```
x = ...
x = layers.Conv2D(32, 3, use_bias=False)(x)
x = layers.BatchNormalization()(x)
```


Because the output of the Conv2D layer gets normalized, the layer doesn't need its own bias vector.

Parameters : [gamma weights, beta weights, moving_mean (non-trainable), moving_variance(non-trainable)] = 32 x 4 = 128

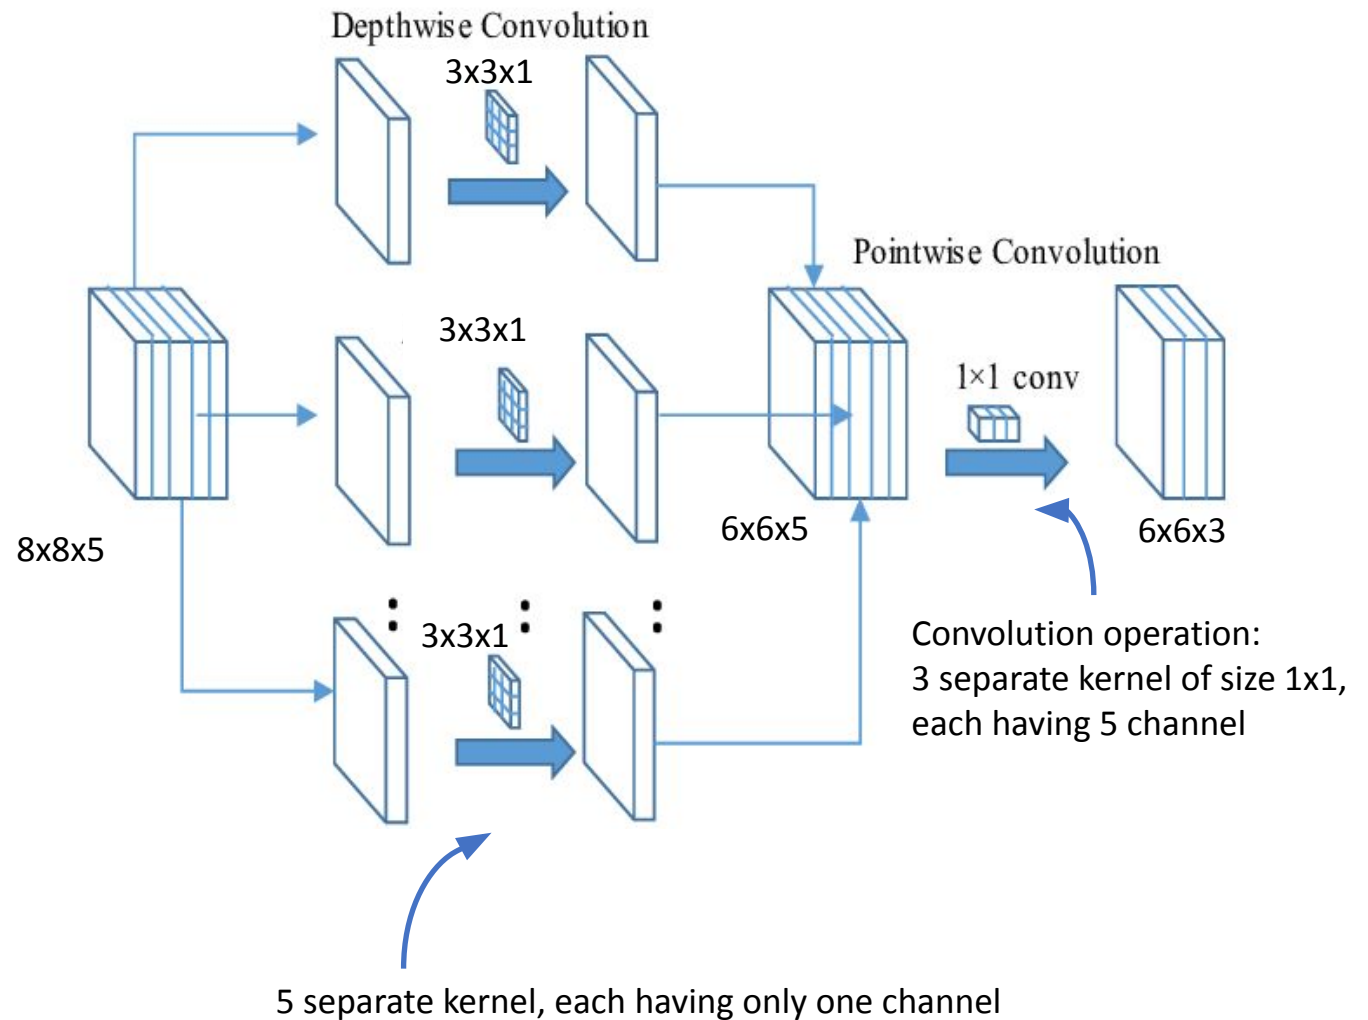
Depthwise separable convolution

Depthwise separable convolutions = Depthwise Conv. + Pointwise Conv.

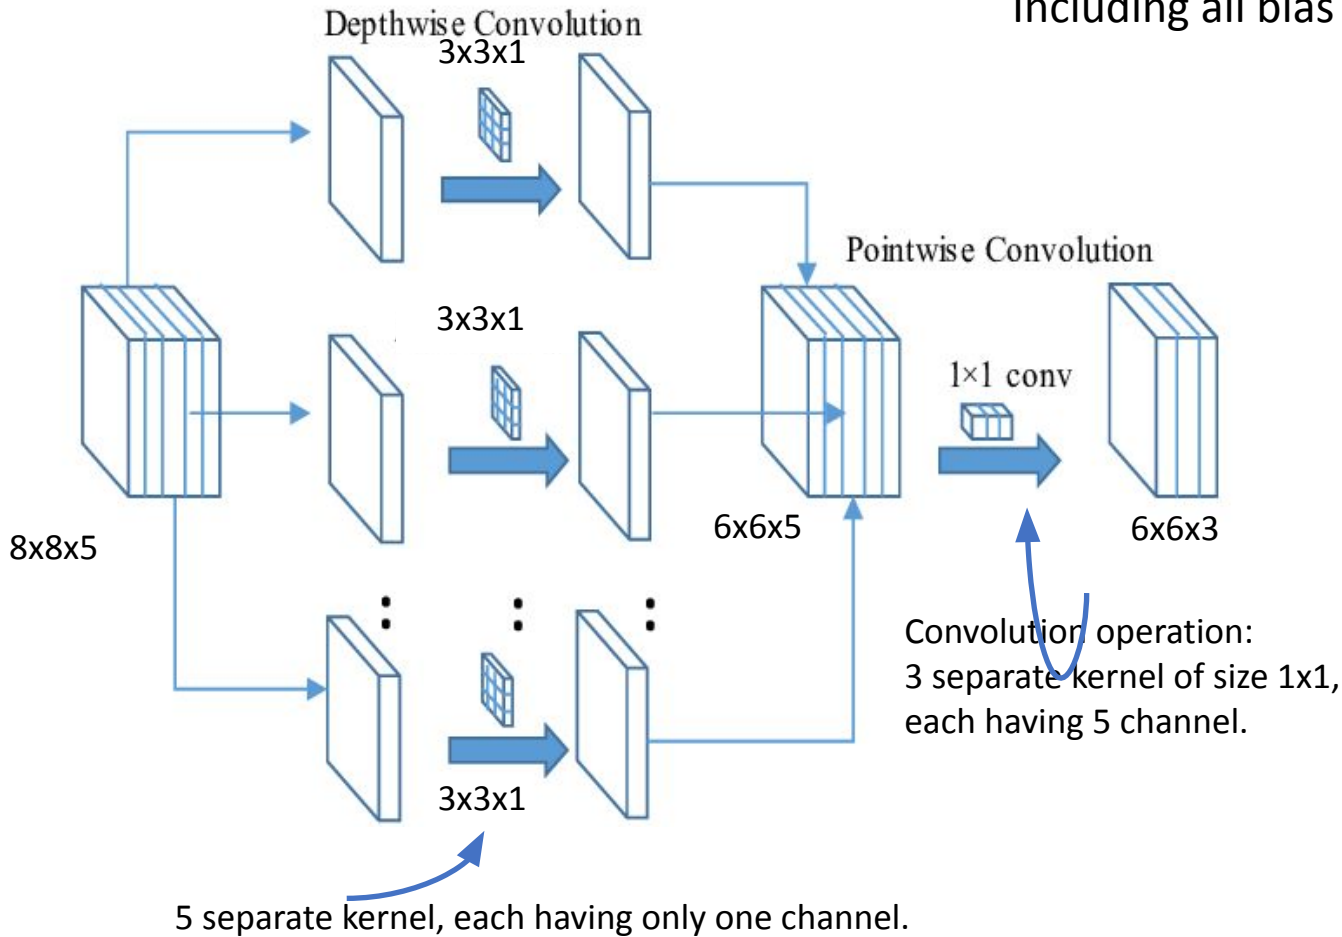
Revisiting Convolution Operation



$$\text{Output shape after Conv} = \frac{n+2p-f}{s} + 1$$



Depthwise separable convolution



Parameter Calculation : $5 \times (3 \times 3) + 5 + (5 \times 1 \times 1 \times 3) + 3 = 68$
Including all bias

Depthwise C.

Pointwise C.

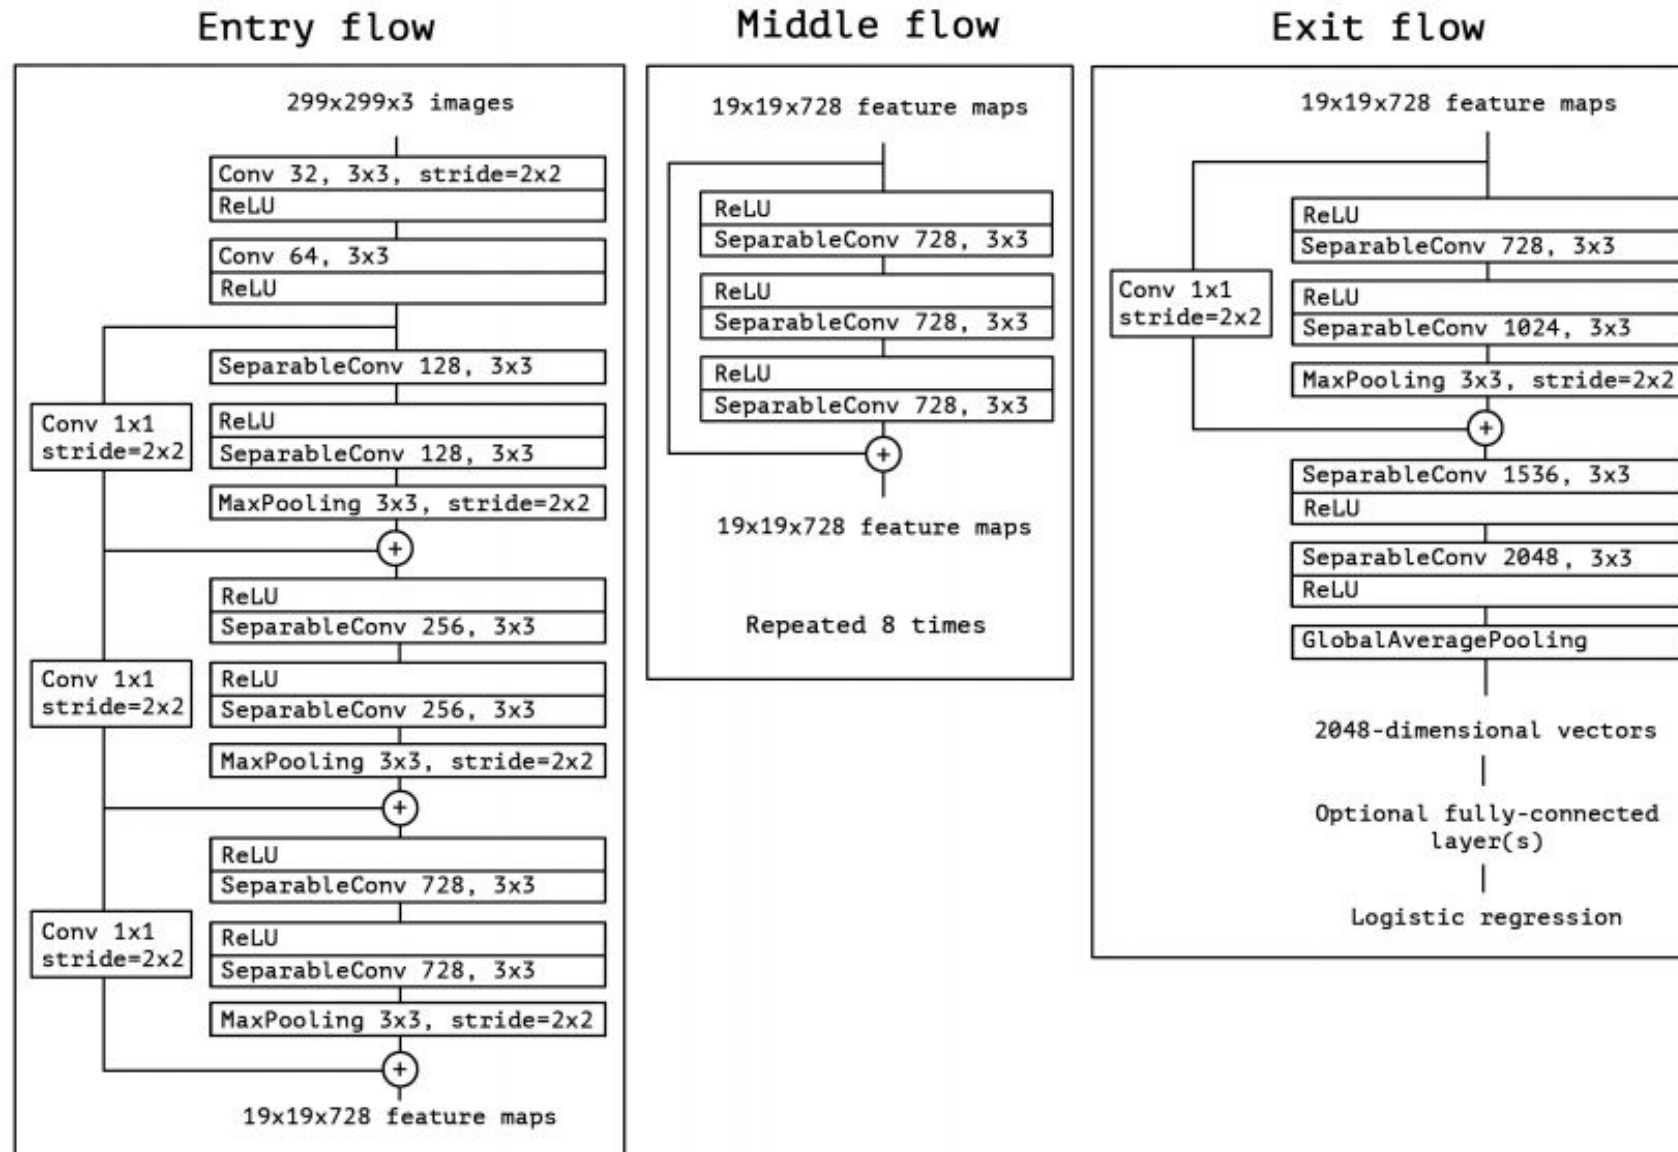
Depthwise C. = N. of in channel x (f x f) + N. of in channel

Pointwise C. = N. of in channel x N. of filter + N. of filter

Keras { `x = layers.SeparableConv2D(64, 3, activation="relu", padding="same") (x)`

In keras, the bias for Depthwise C. is not implemented.

Xception Model



Creating a Mini-Xception Model (a smaller version)

Don't forget input rescaling!

```
inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)

x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=5, use_bias=False)(x)
```

We use the same data augmentation configuration as before.

```
for size in [32, 64, 128, 256, 512]:
    residual = x

    x = layers.BatchNormalization()(x)
    x = layers.Activation("relu")(x)
    x = layers.SeparableConv2D(size, 3, padding="same", use_bias=False)(x)

    x = layers.BatchNormalization()(x)
    x = layers.Activation("relu")(x)
    x = layers.SeparableConv2D(size, 3, padding="same", use_bias=False)(x)

    x = layers.MaxPooling2D(3, strides=2, padding="same")(x)

    residual = layers.Conv2D(
        size, 1, strides=2, padding="same", use_bias=False)(residual)
    x = layers.add([x, residual])

x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

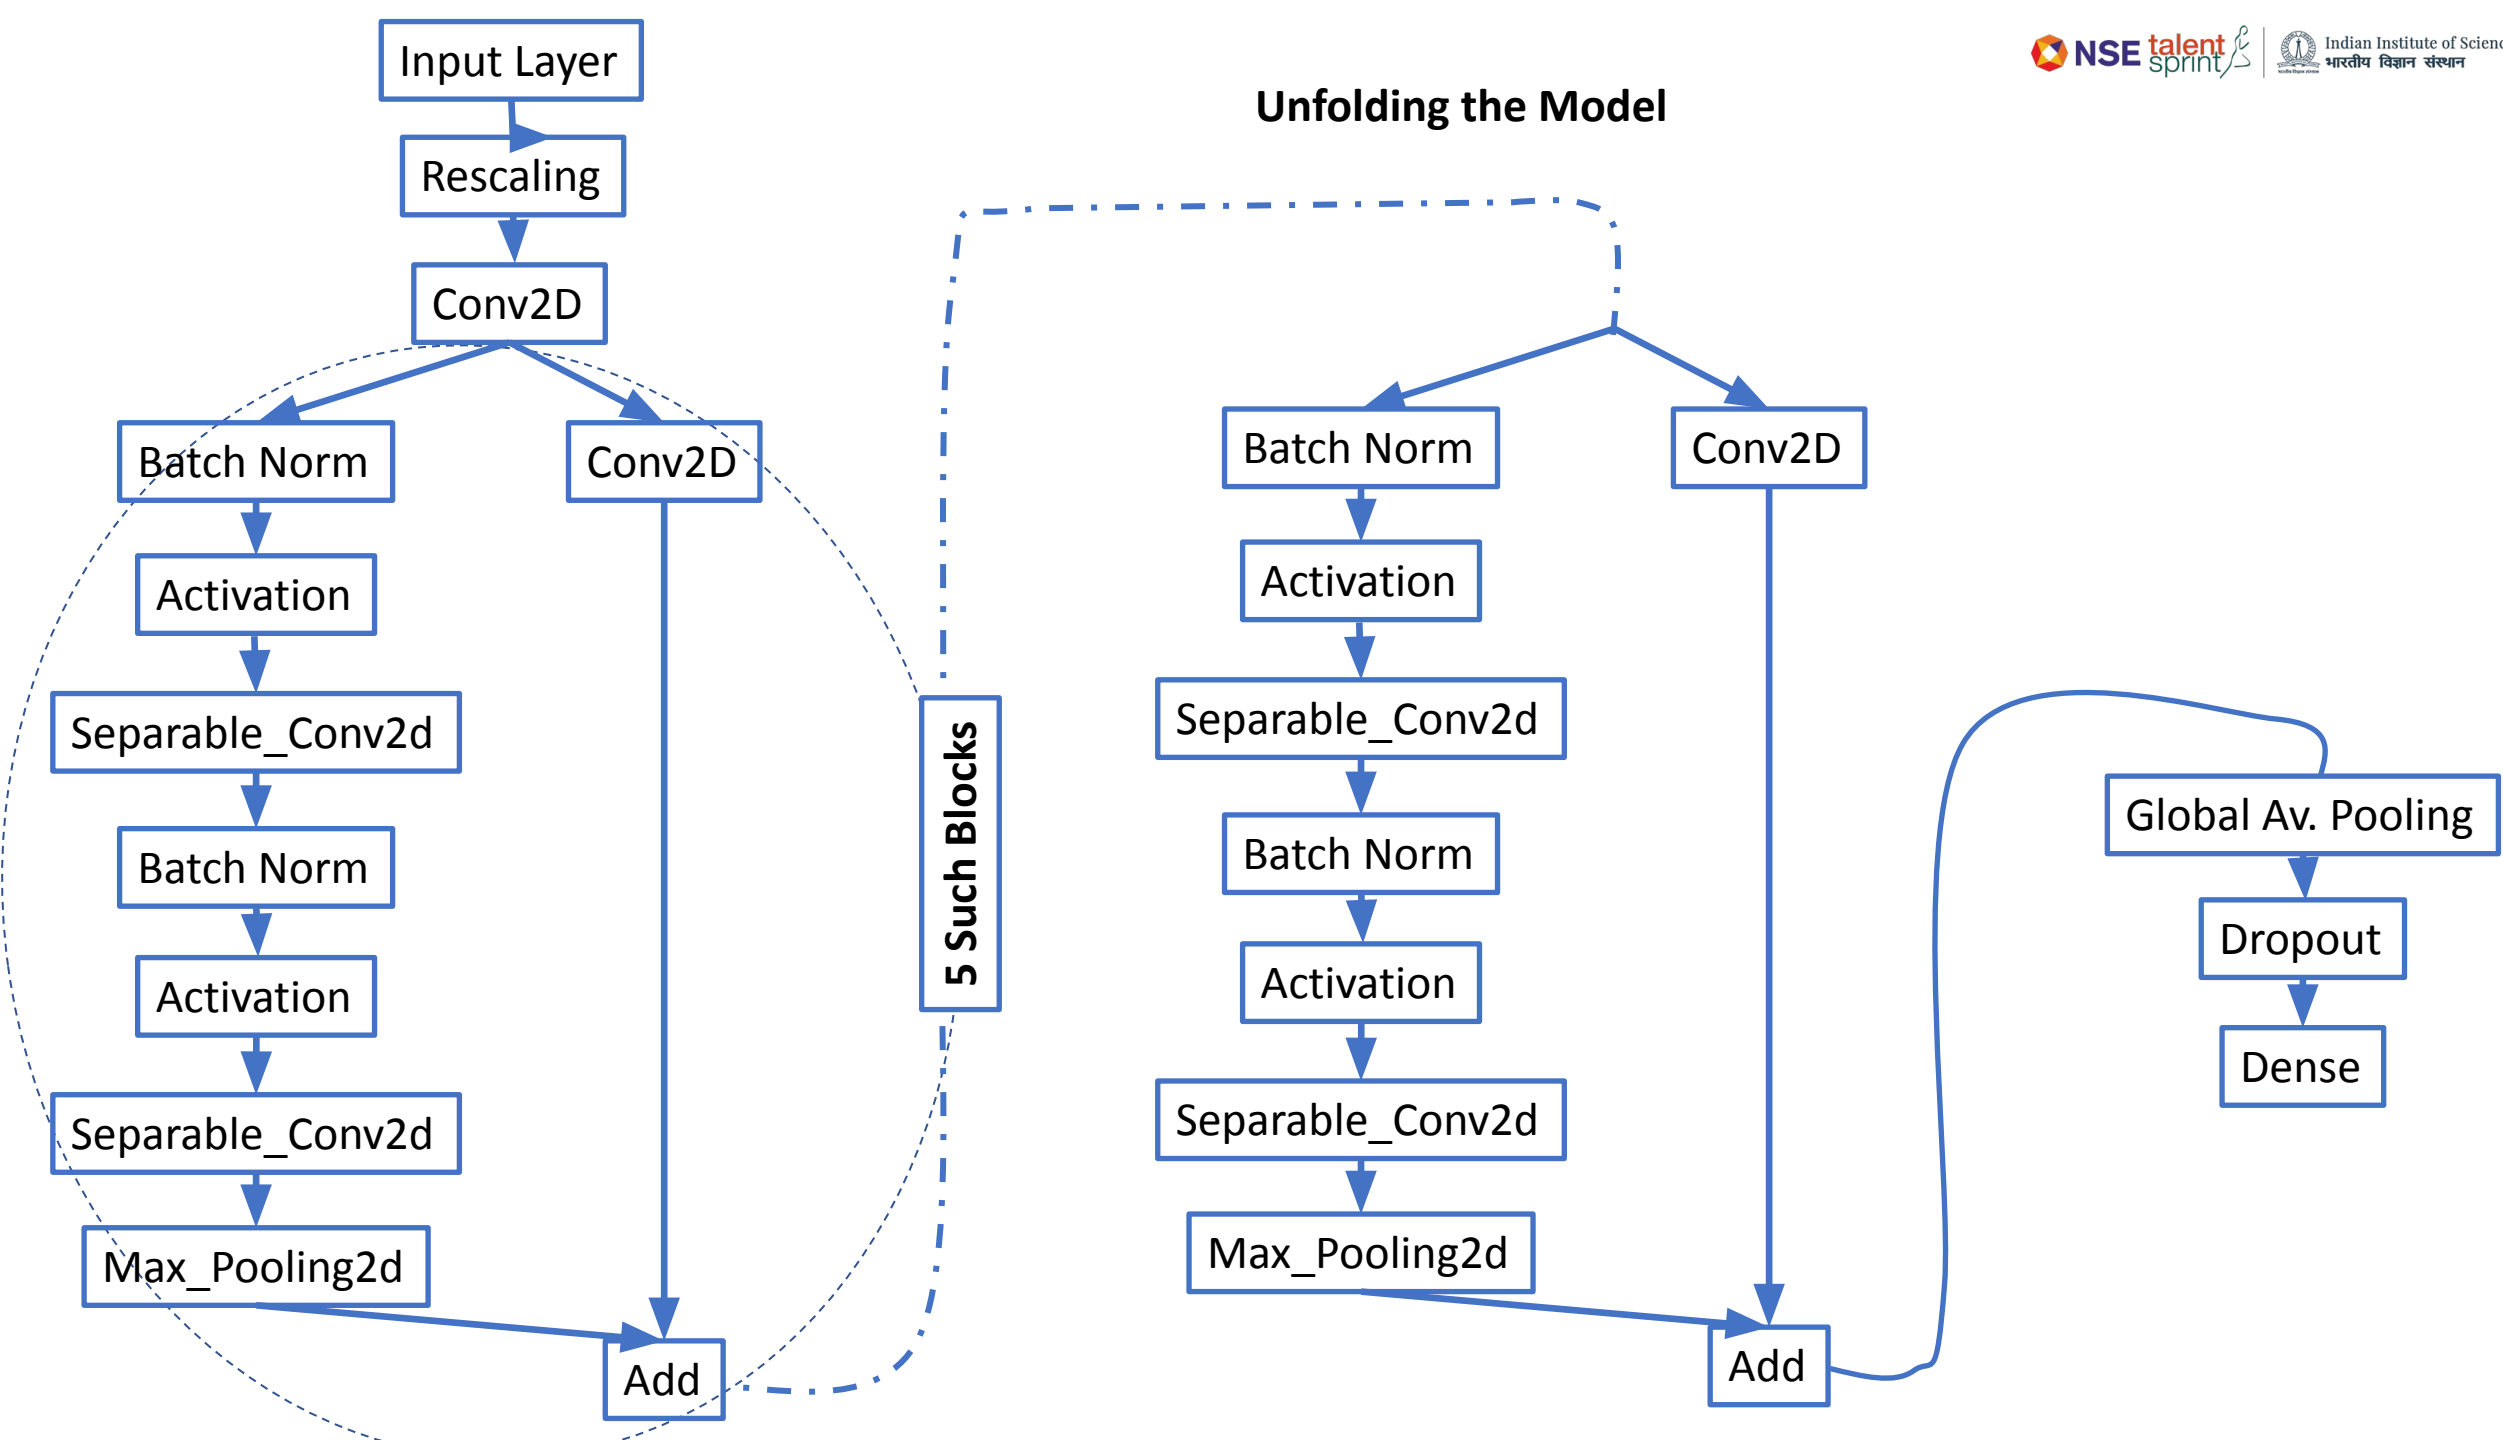
In the original model, we used a Flatten layer before the Dense layer. Here, we go with a GlobalAveragePooling2D layer.

Like in the original model, we add a dropout layer for regularization.

We apply a series of convolutional blocks with increasing feature depth. Each block consists of two batch-normalized depthwise separable convolution layers and a max pooling layer, with a residual connection around the entire block.

Note that the assumption that underlies separable convolution, “feature channels are largely independent,” does not hold for RGB images! Red, green, and blue color channels are actually highly correlated in natural images. As such, the first layer in our model is a regular Conv2D layer. We'll start using SeparableConv2D afterwards.

Unfolding the Model



PART : B

Interpret what CNN's learn

1. Visualizing intermediate activations

```
path = '/content/convnet_from_scratch_with_augmentation.keras'  
model = keras.models.load_model(path)
```

```
plot_model(model)
```

```
from tensorflow.keras import layers
```

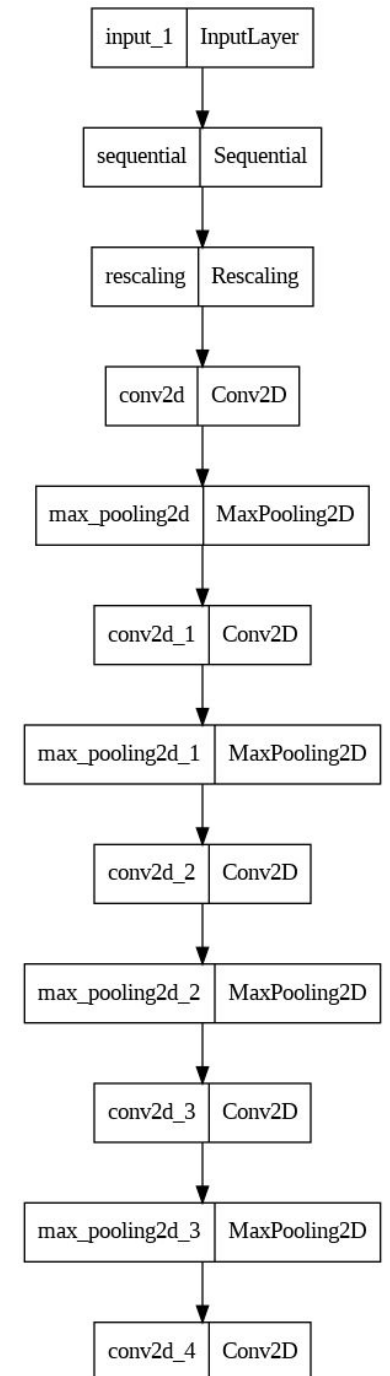
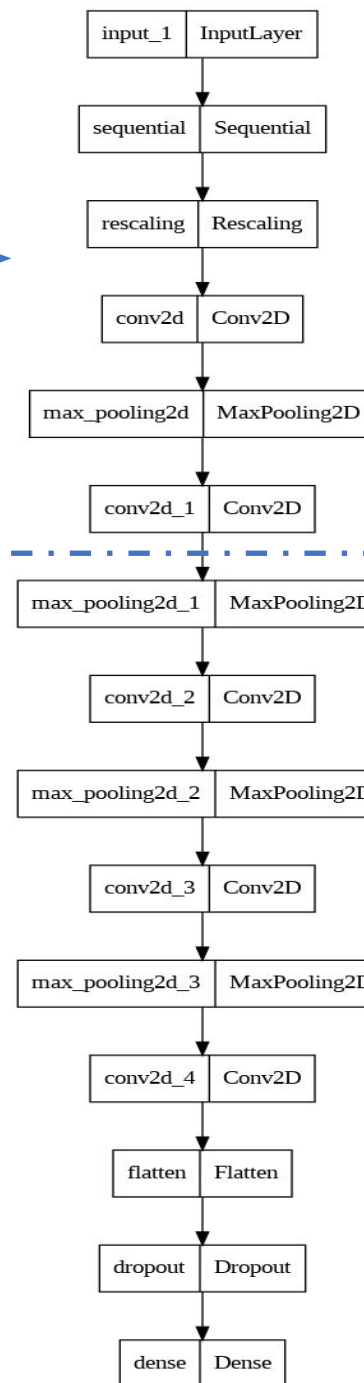
```
layer_outputs = []  
layer_names = []
```

```
for layer in model.layers:  
    if isinstance(layer, (layers.Conv2D, layers.MaxPooling2D)):  
        layer_outputs.append(layer.output)  
        layer_names.append(layer.name)  
activation_model = keras.Model(inputs=model.input, outputs=layer_outputs)  
plot_model(activation_model)
```

Pre-processing the image and feeding into the model :

```
activations = activation_model.predict(img_tensor)  
print(f"No. of outputs= {len(activations)}")  
  
first_layer_feature_maps = activations[0]  
print(f"first_layer_activation.shape= {first_layer_feature_maps.shape}")
```

```
1/1 [=====] - 0s 420ms/step  
No. of outputs= 9  
first_layer_activation.shape= (1, 178, 178, 32)
```



layer_outputs

```
[<KerasTensor: shape=(None, 178, 178, 32) dtype=float32 (created by layer 'conv2d')>,
<KerasTensor: shape=(None, 89, 89, 32) dtype=float32 (created by layer 'max_pooling2d')>,
<KerasTensor: shape=(None, 87, 87, 64) dtype=float32 (created by layer 'conv2d_1')>,
<KerasTensor: shape=(None, 43, 43, 64) dtype=float32 (created by layer 'max_pooling2d_1')>,
<KerasTensor: shape=(None, 41, 41, 128) dtype=float32 (created by layer 'conv2d_2')>,
<KerasTensor: shape=(None, 20, 20, 128) dtype=float32 (created by layer 'max_pooling2d_2')>,
<KerasTensor: shape=(None, 18, 18, 256) dtype=float32 (created by layer 'conv2d_3')>,
<KerasTensor: shape=(None, 9, 9, 256) dtype=float32 (created by layer 'max_pooling2d_3')>,
<KerasTensor: shape=(None, 7, 7, 256) dtype=float32 (created by layer 'conv2d_4')>]
```

layer_names

```
['conv2d',
'max_pooling2d',
'conv2d_1',
'max_pooling2d_1',
'conv2d_2',
'max_pooling2d_2',
'conv2d_3',
'max_pooling2d_3',
'conv2d_4']
```

```
activations = activation_model.predict(img_tensor)
print(f"No. of outputs= {len(activations)}")
```

```
first_layer_feature_maps = activations[0]
print(f"first_layer_activation.shape= {first_layer_feature_maps.shape}")
```

```
1/1 [=====] - 0s 420ms/step
No. of outputs= 9
first_layer_activation.shape= (1, 178, 178, 32)
```

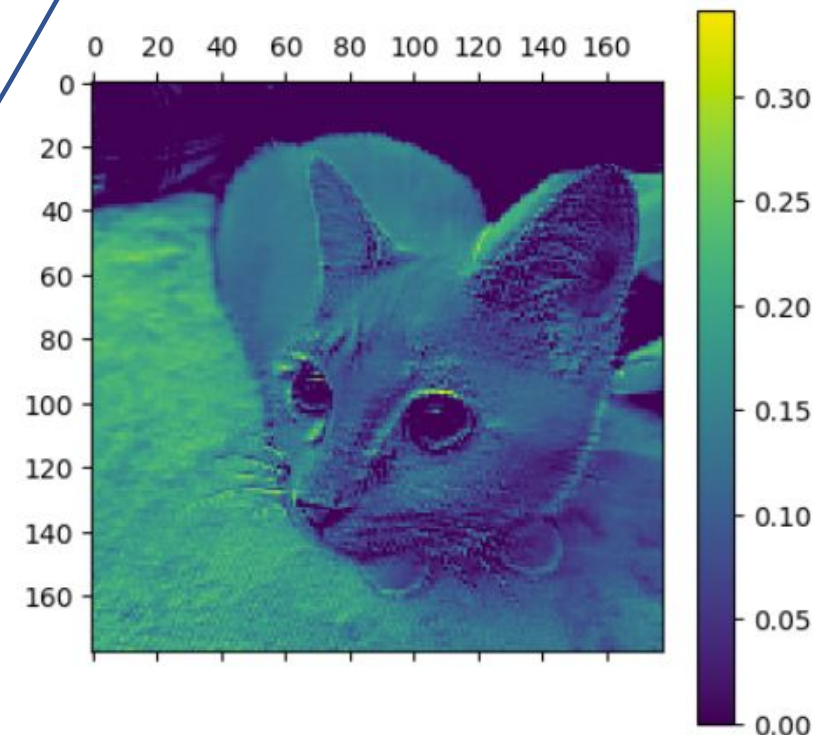
```
print(type(activations))
print(len(activations))
activations[0].shape
```

```
<class 'list'>
9
(1, 178, 178, 32)
```

```
activations[0][0,:,:,:0].shape
(178, 178)
```

```
import matplotlib.pyplot as plt
plt.matshow(first_layer_feature_maps[0, :, :, 0], cmap="viridis")
plt.colorbar()
```

<matplotlib.colorbar.Colorbar at 0x7f38bc463280>




[Colormaps link](#)

conv2d (Conv2D)

(None, 178, 178, 32)

conv2d




layer_names

```
['conv2d',  
'max_pooling2d',  
'conv2d_1',  
'max_pooling2d_1',  
'conv2d_2',  
'max_pooling2d_2',  
'conv2d_3',  
'max_pooling2d_3',  
'conv2d_4']
```

16x16=256

conv2d_4



conv2d_4 (Conv2D)

(None, 7, 7, 256)

Findings :

- The first layer acts as a collection of various edge detectors.
- As we go deeper, the activations become increasingly abstract and less visually interpretable. They begin to encode higher-level concepts.
- The sparsity of the activations increases with the depth of the layer: In the first layer, almost all filters are activated by the input image, but in the following layers, more and more filters are blank. This means the pattern encoded by the filter isn't found in the input image

2. Visualising ConvNet filters

- Ask the question: What kind of an input image will excite the filter?
- What should the input image be so that we see a (yellow) feature map?

```
model = keras.applications.xception.Xception(weights="imagenet", include_top=False)
```

```
for layer in model.layers:
    if isinstance(layer, (keras.layers.Conv2D, keras.layers.SeparableConv2D)):
        print(layer.name)
```

```
block1_conv1
block1_conv2
block2_sepconv1
block2_sepconv2
conv2d_4
block3_sepconv1
block3_sepconv2
```

```
# Creating a feature extractor model
layer_name = "block3_sepconv1"
layer = model.get_layer(name=layer_name)
feature_extractor = keras.Model(inputs=model.input, outputs=layer.output)
feature_extractor.summary()
```

input_2 (InputLayer)	[(None, None, None, 3)]	0	0
block1_conv1 (Conv2D)	(None, None, None, 32)	864	['input_2[0][0]']
block1_conv1_bn (BatchNormalization)	(None, None, None, 32)	128	['block1_conv1[0][0]']
block1_conv1_act (Activation)	(None, None, None, 32)	0	['block1_conv1_bn[0][0]']
block1_conv2 (Conv2D)	(None, None, None, 64)	18432	['block1_conv1_act[0][0]']
block1_conv2_bn (BatchNormalization)	(None, None, None, 64)	256	['block1_conv2[0][0]']
block1_conv2_act (Activation)	(None, None, None, 64)	0	['block1_conv2_bn[0][0]']
block2_sepconv1 (SeparableConv2D)	(None, None, None, 128)	8768	['block1_conv2_act[0][0]']
block2_sepconv1_bn (BatchNormalization)	(None, None, None, 128)	512	['block2_sepconv1[0][0]']
block2_sepconv2_act (Activation)	(None, None, None, 128)	0	['block2_sepconv1_bn[0][0]']
block2_sepconv2 (SeparableConv2D)	(None, None, None, 128)	17536	['block2_sepconv2_act[0][0]']
block2_sepconv2_bn (BatchNormalization)	(None, None, None, 128)	512	['block2_sepconv2[0][0]']
conv2d_4 (Conv2D)	(None, None, None, 128)	8192	['block1_conv2_act[0][0]']
block2_pool (MaxPooling2D)	(None, None, None, 128)	0	['block2_sepconv2_bn[0][0]']
batch_normalization_4 (BatchNormalization)	(None, None, None, 128)	512	['conv2d_4[0][0]']
add_12 (Add)	(None, None, None, 128)	0	['block2_pool[0][0]', 'batch_normalization_4[0][0]']
block3_sepconv1_act (Activation)	(None, None, None, 128)	0	['add_12[0][0]']
block3_sepconv1 (SeparableConv2D)	(None, None, None, 256)	33920	['block3_sepconv1_act[0][0]']



Problem formulation:

We want that image as an input to the filter which excites the filter most. What does it mean?

Different input images are passed through the filter and the mean of all the pixel values is observed in output from the filter.

Initialize with a random image and calculate the mean of the activation values from the filter (loss function).

Calculate gradient of loss wrt to image. Apply gradient ascent. Get the updated image. Iterate...



```
import tensorflow as tf

def compute_loss(image, filter_index):
    activation = feature_extractor(image)
    filter_activation = activation[:, 2:-2, 2:-2, filter_index]
    return tf.reduce_mean(filter_activation)
```

Note that we avoid border artifacts by only involving non-border pixels in the loss; we discard the first two pixels along the sides of the activation.

```
@tf.function
def gradient_ascent_step(image, filter_index, learning_rate):
    with tf.GradientTape() as tape:
        tape.watch(image)
        loss = compute_loss(image, filter_index)
    grads = tape.gradient(loss, image)
    grads = tf.math.l2_normalize(grads)
    image += learning_rate * grads
    return image
```

```
img_width = 200
img_height = 200

def generate_filter_pattern(filter_index):
    iterations = 30
    learning_rate = 10.
    image = tf.random.uniform(minval=0.4, maxval=0.6, shape=(1, img_width, img_height, 3))
    for i in range(iterations):
        image = gradient_ascent_step(image, filter_index, learning_rate)
    return image[0].numpy()
```

```
def deprocess_image(image):
    image -= image.mean()
    image /= image.std()
    image *= 64
    image += 128
    image = np.clip(image, 0, 255).astype("uint8")
    image = image[25:-25, 25:-25, :]
    return image

plt.axis("off")
plt.imshow(deprocess_image(generate_filter_pattern(filter_index=2)))
```


Thanks !